

# HW6

Jon Griffith and Annabelle Cunningham

2025-04-21

```
library(ISLR2)

## Warning: package 'ISLR2' was built under R version 4.4.3
library(glmnet)

## Warning: package 'glmnet' was built under R version 4.4.3
## Loading required package: Matrix
## Loaded glmnet 4.1-8
library(leaps)

## Warning: package 'leaps' was built under R version 4.4.3
library(pls)

## Warning: package 'pls' was built under R version 4.4.3
##
## Attaching package: 'pls'
## The following object is masked from 'package:stats':
##      loadings
```

8

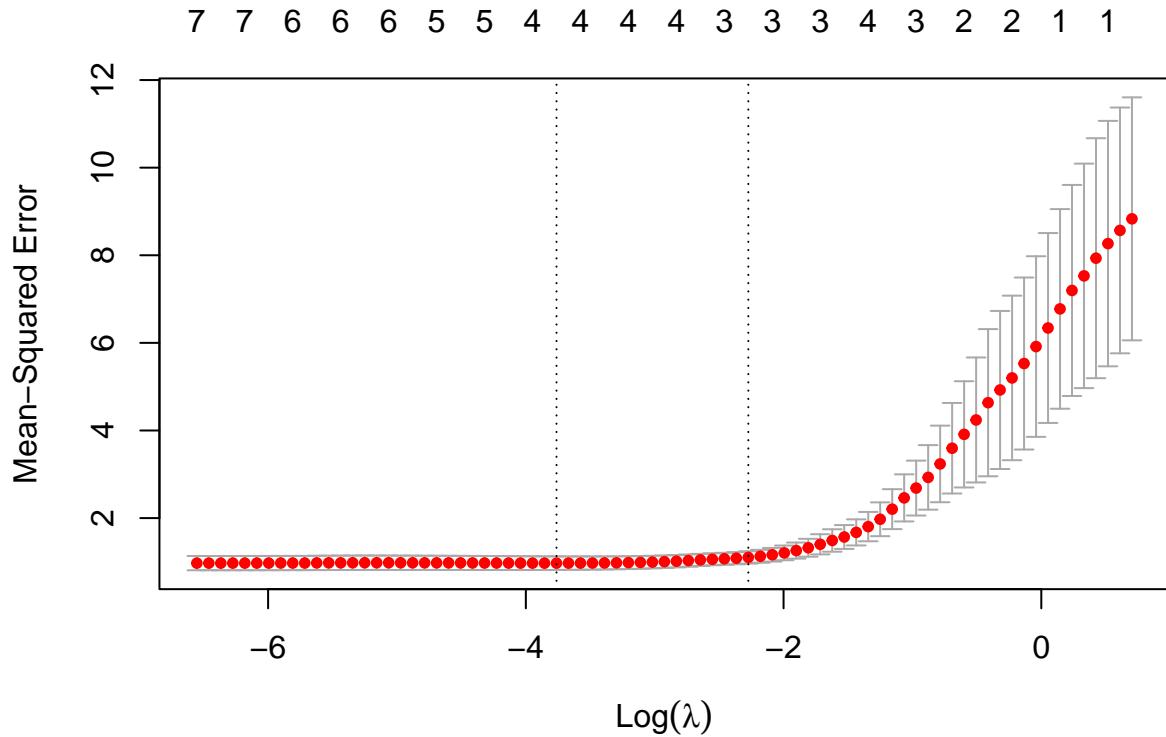
(e)

```
set.seed(1)
# Generate data
x <- rnorm(100)
e <- rnorm(100)
Y <- 4 - 3*x - 2*x^2 + 1.25*x^3 + e

df <- data.frame(x,Y)

set.seed(1)
cv.out <- cv.glmnet(poly(df$x,10,raw=TRUE), df$Y, alpha=1)

plot(cv.out)
```



```
lasso.mod <- glmnet(poly(df$x, 10, raw=TRUE), df$Y, alpha=1, lambda=cv.out$lambda.min)
predict(lasso.mod, type='coefficients', s=cv.out$lambda.min)
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##           s1
## (Intercept) 4.0446480
## 1          -2.3121547
## 2          -2.1331883
## 3           0.5836176
## 4            .
## 5           0.1110500
## 6            .
## 7            .
## 8            .
## 9            .
## 10           .
```

Our resulting coefficient predictions based on the optimal lambda from CV come close to the true values for  $\beta_0$ ,  $\beta_1$ , and  $\beta_2$ , while  $\beta_3$  is about half of its actual value and  $\beta_5$  is included even though there was no  $X^5$  in the true equation. It still did a pretty good job of selecting close to the actual number of variables in the model and was only off by one.

(f)

```
set.seed(1)
df$Y <- 4 - 1.25*x^7 + e

bestsub.mod <- regsubsets(Y ~ poly(x, 10, raw=TRUE), data=df)

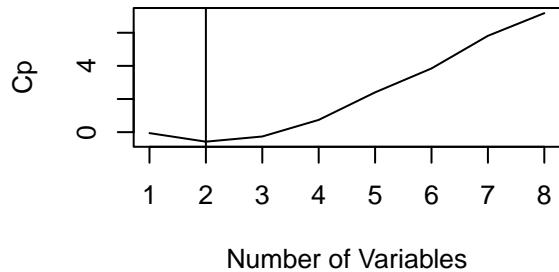
reg.summary <- summary(bestsub.mod)
```

```

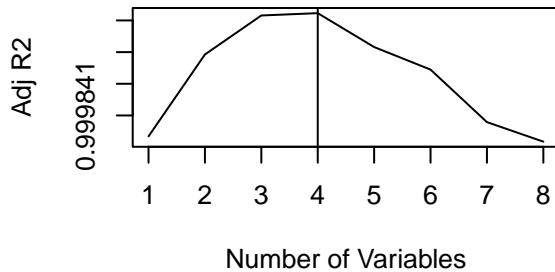
par(mfrow=c(2,2))
# PLOT of cp vs num variables
plot(reg.summary$cp,
  xlab='Number of Variables',
  ylab='Cp',
  type='l',
  main= paste('Optimal:', which.min(reg.summary$cp)))
abline(v=which.min(reg.summary$cp))
# PLOT of adjr2 vs num variables
plot(reg.summary$adjr2,
  xlab='Number of Variables',
  ylab='Adj R2',
  type='l',
  main= paste('Optimal:', which.max(reg.summary$adjr2)))
abline(v=which.max(reg.summary$adjr2))
# PLOT of bic vs num variables
plot(reg.summary$bic,
  xlab='Number of Variables',
  ylab='BIC',
  type='l',
  main= paste('Optimal:', which.min(reg.summary$bic)))
abline(v=which.min(reg.summary$bic))

```

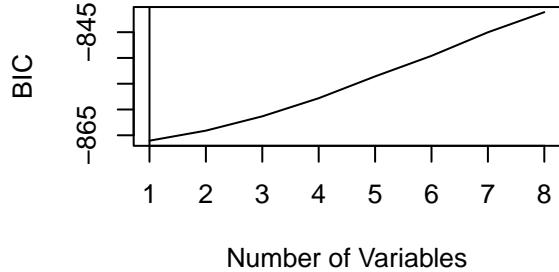
**Optimal: 2**



**Optimal: 4**



**Optimal: 1**



```
coef(bestsub.mod, 1)
```

```

##           (Intercept) poly(x, 10, raw = TRUE)7
##            3.95894          -1.24923

```

```
coef(bestsub.mod, 2)
```

```

##          (Intercept) poly(x, 10, raw = TRUE)2 poly(x, 10, raw = TRUE)7
##          4.0704904           -0.1417084            -1.2484448
coef(bestsub.mod, 4)

##          (Intercept) poly(x, 10, raw = TRUE)1 poly(x, 10, raw = TRUE)2
##          4.0762524           0.2914016           -0.1617671
## poly(x, 10, raw = TRUE)3 poly(x, 10, raw = TRUE)7
##          -0.2526527           -1.2408662
coef(bestsub.mod, 1)

##          (Intercept) poly(x, 10, raw = TRUE)7
##          3.95894           -1.24923
coef(bestsub.mod, 2)

##          (Intercept) poly(x, 10, raw = TRUE)2 poly(x, 10, raw = TRUE)7
##          4.0704904           -0.1417084            -1.2484448
coef(bestsub.mod, 4)

##          (Intercept) poly(x, 10, raw = TRUE)1 poly(x, 10, raw = TRUE)2
##          4.0762524           0.2914016           -0.1617671
## poly(x, 10, raw = TRUE)3 poly(x, 10, raw = TRUE)7
##          -0.2526527           -1.2408662

```

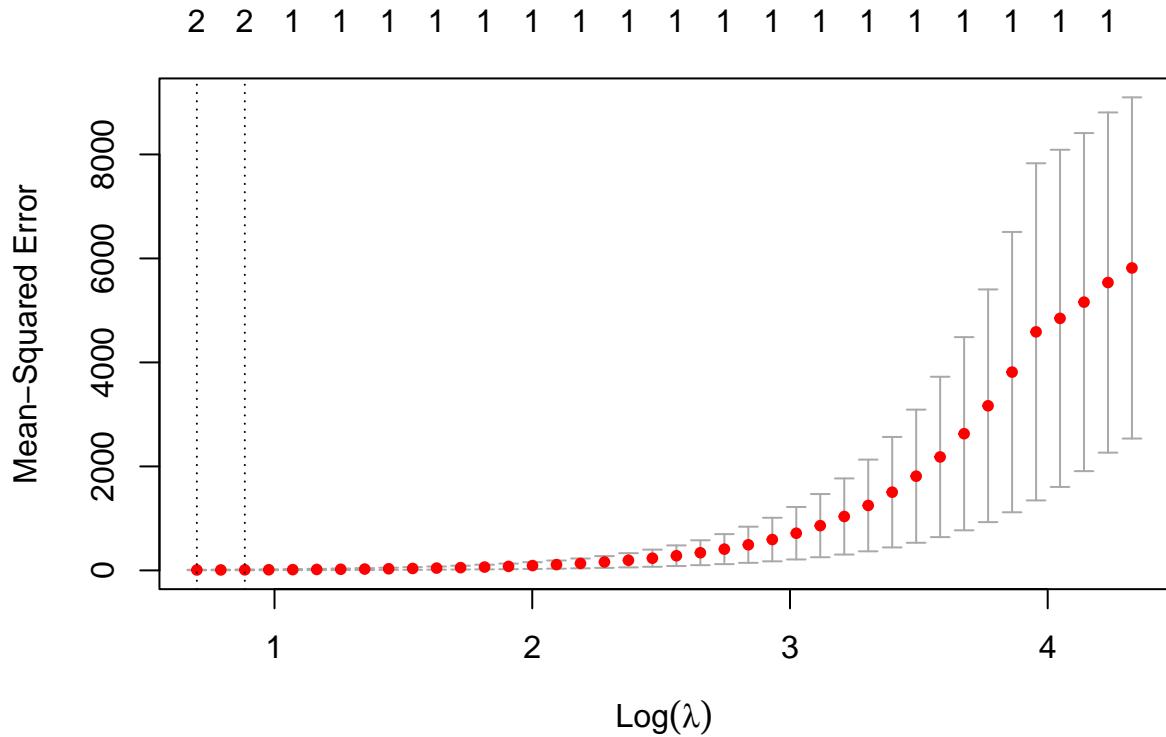
The best models as measured by BIC, CP, and Adj  $R^2$  determined best subsets of  $(X^7)$ ,  $(X^2, X^7)$ , and  $(X, X^2, X^3, \text{ and } X^7)$ , respectively. Each model included the predictor  $X^7$  that corresponded to the true predictor, and the intercepts and coefficient for  $X^7$  very closely approximated the true values for each. BIC did the best job but CP and Adj  $R^2$  were nearly as good, adding only negligible effects of extra predictors.

```

set.seed(1)
cv.out <- cv.glmnet(poly(df$x, 10, raw=TRUE), df$Y, alpha=1)

plot(cv.out)

```



```
lasso.mod <- glmnet(poly(df$x, 10, raw=TRUE), df$Y, alpha=1, lambda=cv.out$lambda.min)
predict(lasso.mod, type='coefficients', s=cv.out$lambda.min)
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##                               s1
## (Intercept) 3.81888390
## 1           .
## 2           .
## 3           .
## 4           .
## 5      -0.01814104
## 6           .
## 7     -1.21252256
## 8           .
## 9           .
## 10          .
```

The optimal lambda and corresponding model for Lasso, per CV, includes an intercept value of 3.81 which very closely approximates the true intercept of 4, the  $X^5$  predictor with a coefficient of -0.018, which is not in the true model but is negligible, and the  $X^7$  predictor with a coefficient of -1.212, which very closely approximates the true coefficient of -1.25. The Lasso does a good job of replicating the true model and is about on par with the best subsets method.

## 9

(a)

```
df <- College
head(df)

##                                     Private Apps Accept Enroll Top10perc Top25perc
```

```

## Abilene Christian University Yes 1660 1232 721 23 52
## Adelphi University Yes 2186 1924 512 16 29
## Adrian College Yes 1428 1097 336 22 50
## Agnes Scott College Yes 417 349 137 60 89
## Alaska Pacific University Yes 193 146 55 16 44
## Albertson College Yes 587 479 158 38 62
## F.Undergrad P.Undergrad Outstate Room.Board Books
## Abilene Christian University 2885 537 7440 3300 450
## Adelphi University 2683 1227 12280 6450 750
## Adrian College 1036 99 11250 3750 400
## Agnes Scott College 510 63 12960 5450 450
## Alaska Pacific University 249 869 7560 4120 800
## Albertson College 678 41 13500 3335 500
## Personal PhD Terminal S.F.Ratio perc.alumni Expend
## Abilene Christian University 2200 70 78 18.1 12 7041
## Adelphi University 1500 29 30 12.2 16 10527
## Adrian College 1165 53 66 12.9 30 8735
## Agnes Scott College 875 92 97 7.7 37 19016
## Alaska Pacific University 1500 76 72 11.9 2 10922
## Albertson College 675 67 73 9.4 11 9727
## Grad.Rate
## Abilene Christian University 60
## Adelphi University 56
## Adrian College 54
## Agnes Scott College 59
## Alaska Pacific University 15
## Albertson College 55

set.seed(1)
train <- sample(1:nrow(df), round(0.7*nrow(df)))

```

(b)

```

ols.mod <- lm(Apps ~ ., data=df[train,])
yhat <- predict(ols.mod, newdata = df[-train,])
MSE <- mean((yhat - df$Apps[-train])^2)
MSE

## [1] 1266407

```

(c)

```

set.seed(1)
train_split <- model.matrix(Apps ~ .-1, data=df[train,])
test_split <- model.matrix(Apps ~ .-1, data=df[-train,])
ridge.cv <- cv.glmnet(train_split, df$Apps[train], alpha=0)
yhat <- predict(ridge.cv, newx = test_split, s = ridge.cv$lambda.min)
MSE_ridge <- mean((yhat - df$Apps[-train])^2)
MSE_ridge

## [1] 1122733
ridge.coef <- predict(ridge.cv, type='coefficients', s=ridge.cv$lambda.min)[1:19,]
ridge.coef

```

```

##   (Intercept) PrivateNo PrivateYes      Accept      Enroll
## -1.835673e+03 2.319012e+02 -2.327310e+02 1.097370e+00 3.870257e-01
##   Top10perc    Top25perc F.Undergrad P.Undergrad     Outstate
##  2.656554e+01 2.809031e-02  5.802492e-02 3.525905e-02 -3.238312e-02
##   Room.Board      Books    Personal       PhD     Terminal
##  2.182992e-01 3.007454e-01 -3.084963e-02 -4.137527e+00 -4.295744e+00
##   S.F.Ratio  perc.alumni     Expend   Grad.Rate
##  1.473171e+01 -4.595790e+00  6.404545e-02 1.003295e+01

```

(d)

```

set.seed(1)
lasso.cv <- cv.glmnet(train_split, df$Apps[train], alpha=1)
yhat <- predict(lasso.cv, newx = test_split, s = lasso.cv$lambda.min)
MSE_lasso <- mean((yhat - df$Apps[-train])^2)
lasso.coef <- predict(lasso.cv, type='coefficients', s=lasso.cv$lambda.min)[1:19,]

MSE_lasso

```

```

## [1] 1259544
lasso.coef

```

```

##   (Intercept) PrivateNo PrivateYes      Accept      Enroll
## -1.067871e+03 5.064555e+02 -1.271310e-09 1.706280e+00 -9.129561e-01
##   Top10perc    Top25perc F.Undergrad P.Undergrad     Outstate
##  5.220892e+01 -1.516972e+01  9.427128e-03 7.078118e-02 -8.550683e-02
##   Room.Board      Books    Personal       PhD     Terminal
##  1.603423e-01 2.622365e-01  0.000000e+00 -9.483740e+00 -8.895288e-02
##   S.F.Ratio  perc.alumni     Expend   Grad.Rate
##  1.530680e+01 1.523357e+00  5.921301e-02 6.948170e+00

```

We have only one predictor zeroed out and several others that are close to zero.

(e)

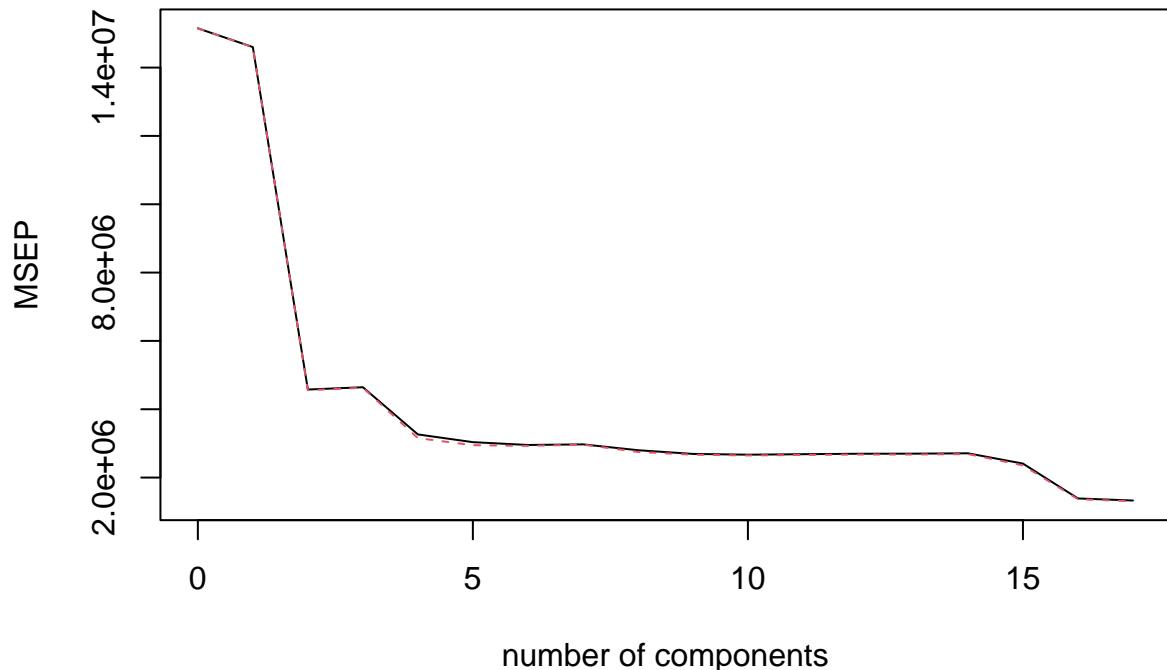
PCR

```

set.seed(1)
pqr.mod <- pcr(Apps ~ ., data=df[train,], scale=TRUE, validation='CV')
validationplot(pqr.mod, val.type='MSEP')

```

## Apps



```
which.min(pcr.mod$validation$PRESS)
```

```
## [1] 17
```

We have a minimum at the model with the same number of components as predictors, so we have  $M = 17 = p$ . This is equivalent to the OLS model, and so we will use this to find our test error equivalent to OLS.

```
M <- 17
yhat <- predict(pcr.mod, df[-train,], ncomp = M)
MSE_pcr <- mean((yhat - df$Apps[-train])^2)

MSE_pcr
```

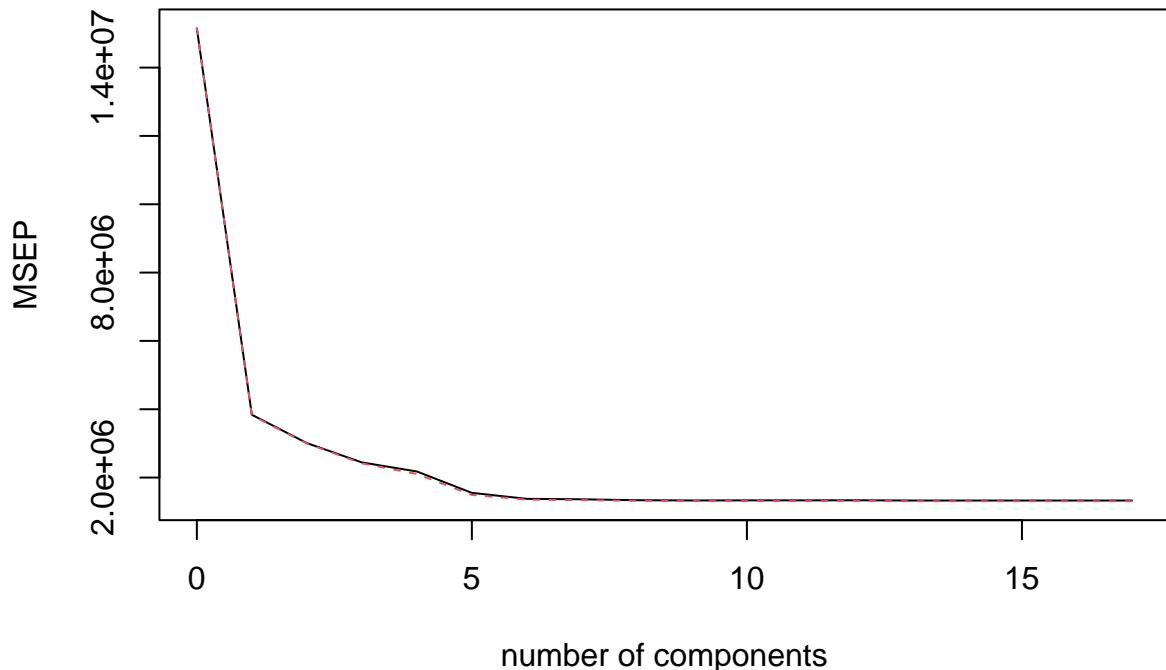
```
## [1] 1266407
```

(f)

PLS

```
set.seed(1)
pls.mod <- plsr(Apps ~ ., data=df[train,], scale=TRUE, validation='CV')
validationplot(pls.mod, val.type='MSEP')
```

## Apps



```
which.min(pls.mod$validation$PRESS)
```

```
## [1] 17
```

We again have a minimum at  $M = 17$ , which is equivalent to the OLS model, and so we will use this to find our test error equivalent to OLS.

```
M <- 17
yhat <- predict(pls.mod, df[-train,], ncomp = M)
MSE_pls <- mean((yhat - df$Apps[-train])^2)

MSE_pls
## [1] 1266407
```

(f)

```
cat('MSE OLS: ', MSE, '\n')
## MSE OLS: 1266407
cat('MSE Ridge: ', MSE_ridge, '\n')
## MSE Ridge: 1122733
cat('MSE Lasso: ', MSE_lasso, '\n')
## MSE Lasso: 1259544
cat('MSE PCR: ', MSE_pcr, '\n')
## MSE PCR: 1266407
```

```
cat('MSE PLS: ', MSE_pls, '\n')
```

```
## MSE PLS: 1266407
```

We achieved our best results from Ridge and Lasso while PCR and PLS just gave the OLS estimate. Out of Ridge and Lasso, Ridge did better.

## 11

(a)

Create train/test splits

```
df <- Boston

set.seed(1)
train <- sample(1:nrow(df), round(0.7*nrow(df)))

train_split <- model.matrix(crim ~ .-1, data=df[train,])
test_split <- model.matrix(crim ~ .-1, data=df[-train,])
```

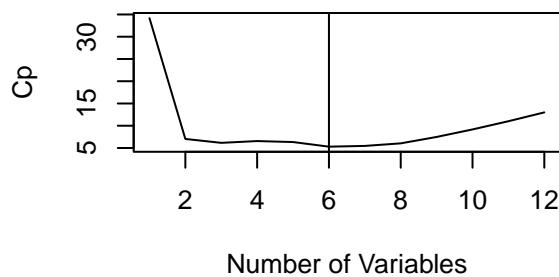
Best Subset Selection

```
bestsub.mod <- regsubsets(crim ~ ., data=df[train,], nvmax=12)

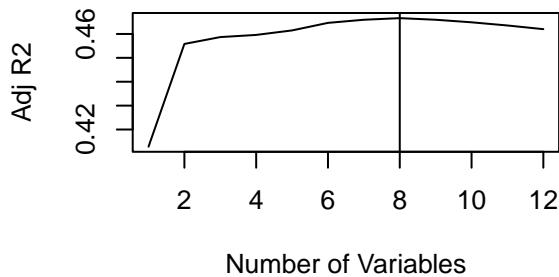
reg.summary <- summary(bestsub.mod)

par(mfrow=c(2,2))
# Plot of cp vs num variables
plot(reg.summary$cp,
  xlab='Number of Variables',
  ylab='Cp',
  type='l',
  main=paste('Optimal:', which.min(reg.summary$cp)))
abline(v=which.min(reg.summary$cp))
# Plot of adjr2 vs num variables
plot(reg.summary$adjr2,
  xlab='Number of Variables',
  ylab='Adj R2',
  type='l',
  main=paste('Optimal:', which.max(reg.summary$adjr2)))
abline(v=which.max(reg.summary$adjr2))
# Plot of bic vs num variables
plot(reg.summary$bic,
  xlab='Number of Variables',
  ylab='BIC',
  type='l',
  main=paste('Optimal:', which.min(reg.summary$bic)))
abline(v=which.min(reg.summary$bic))
```

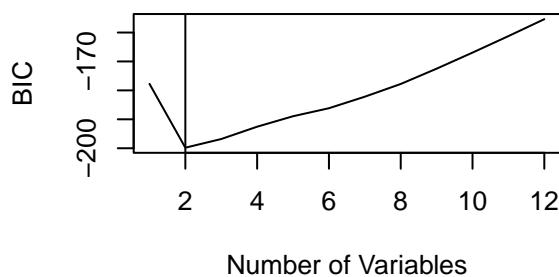
**Optimal: 6**



**Optimal: 8**



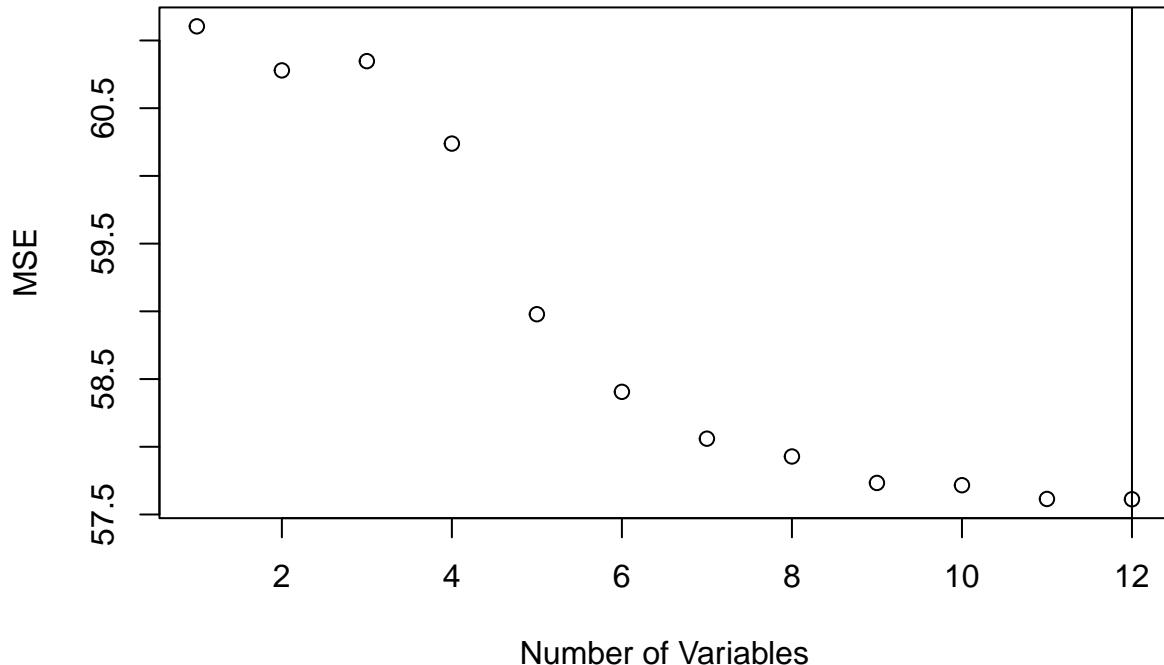
**Optimal: 2**



```
test.mat <- model.matrix(crim ~ ., data=df[-train,])
val.errors <- rep(NA, 12)
for (i in 1:12){
  coefi <- coef(bestsub.mod, id=i)
  pred <- test.mat[,names(coefi)] %*% coefi
  val.errors[i] <- mean((df$crim[-train] - pred)^2)
}
plot(val.errors,
  main='MSE for each model fit\n Optimal: 7',
  xlab='Number of Variables',
  ylab='MSE')
abline(v = which.min(val.errors))
```

### MSE for each model fit

Optimal: 7



```
MSE_bestsub <- val.errors[2]
```

```
MSE_bestsub
```

```
## [1] 60.77902
```

#### Ridge Regression

```
ridge.cv <- cv.glmnet(train_split, df$crim[train], alpha=0)
yhat <- predict(ridge.cv, newx = test_split, s = ridge.cv$lambda.min)
MSE_ridge <- mean((yhat - df$crim[-train])^2)
MSE_ridge
```

```
## [1] 58.75168
```

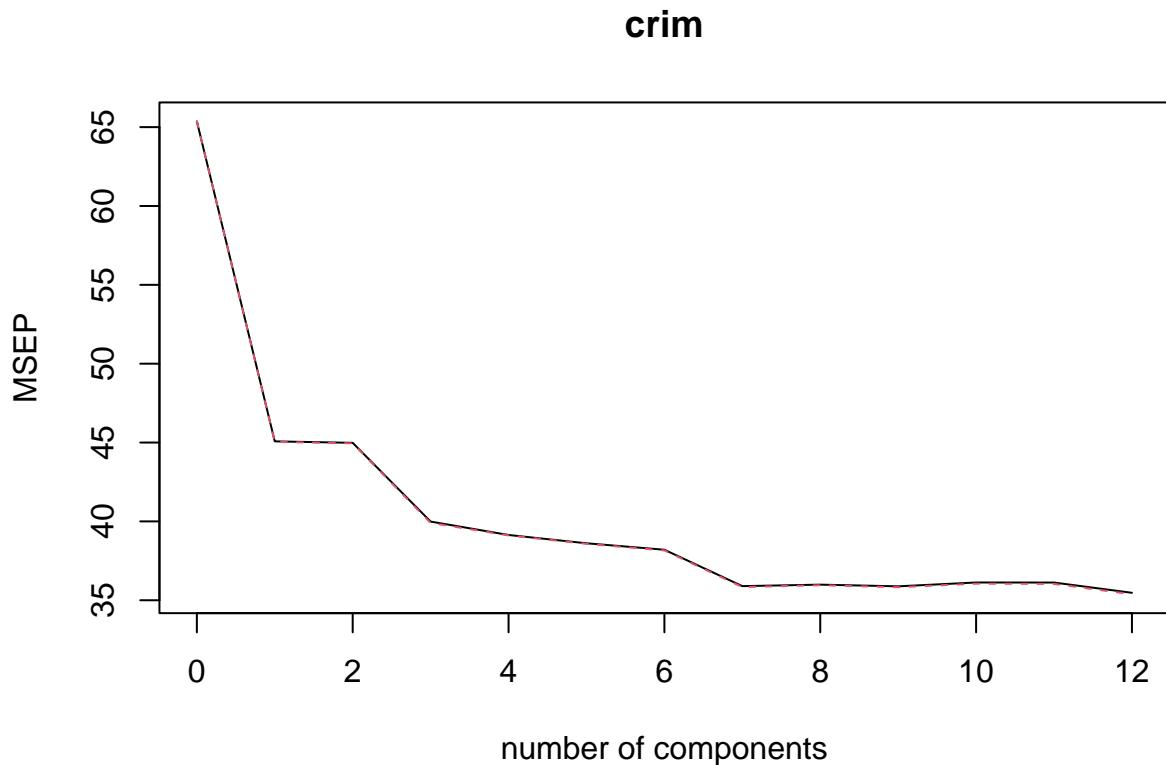
#### Lasso

```
lasso.cv <- cv.glmnet(train_split, df$crim[train], alpha=1)
yhat <- predict(lasso.cv, newx = test_split, s = lasso.cv$lambda.min)
MSE_lasso <- mean((yhat - df$crim[-train])^2)
MSE_lasso
```

```
## [1] 58.06509
```

#### PCR

```
set.seed(1)
pqr.mod <- pqr(crim ~ ., data=df[train,], scale=TRUE, validation='CV')
validationplot(pqr.mod, val.type='MSEP')
```



```
which.min(pcr.mod$validation$PRESS)
```

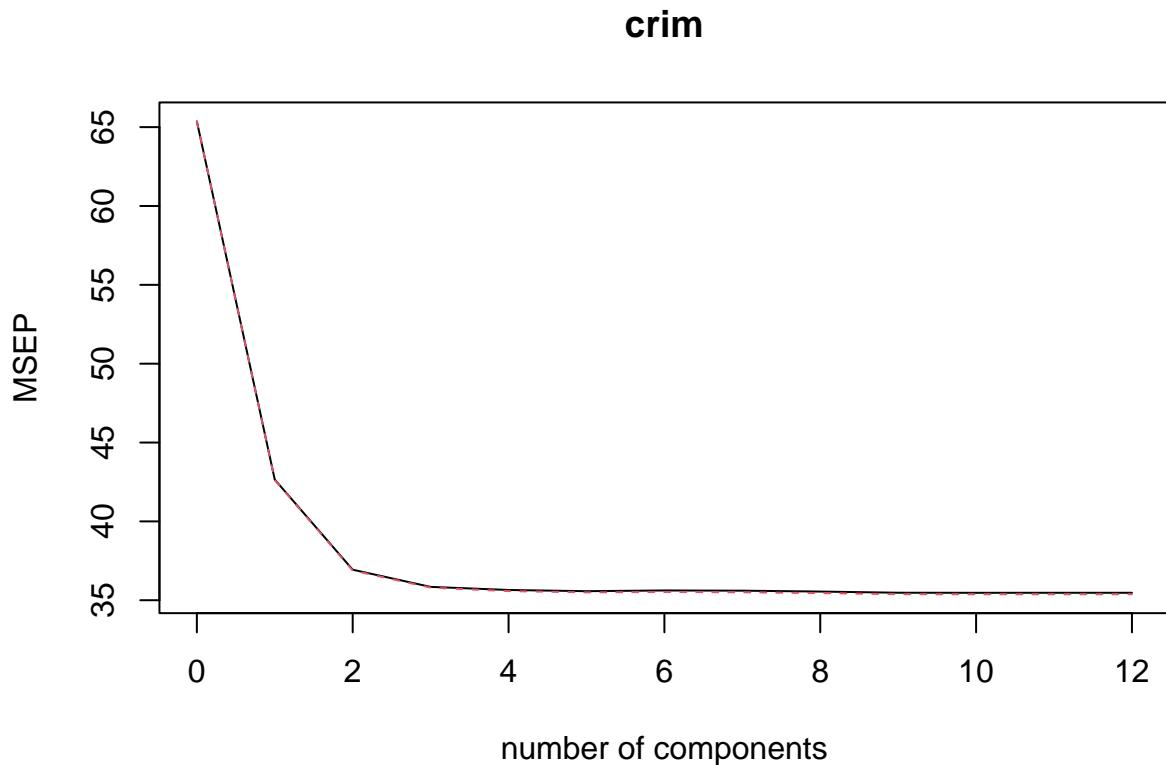
```
## [1] 12
M <- 12
yhat <- predict(pcr.mod, df[-train,], ncomp = M)
MSE_pcr <- mean((yhat - df$crim[-train])^2)
```

```
MSE_pcr
```

```
## [1] 57.61252
```

**PLS**

```
set.seed(1)
pls.mod <- plsr(crim ~ ., data=df[train,], scale=TRUE, validation='CV')
validationplot(pls.mod, val.type='MSEP')
```



```
which.min(pls.mod$validation$PRESS)
```

```
## [1] 10
M <- 10
yhat <- predict(pls.mod, df[-train,], ncomp = M)
MSE_pls <- mean((yhat - df$crim[-train])^2)
```

```
MSE_pls
```

```
## [1] 57.61217
```

(b)

```
cat('MSE Best Subset: ', MSE_bestsub, '\n')
```

```
## MSE Best Subset: 60.77902
```

```
cat('MSE Ridge: ', MSE_ridge, '\n')
```

```
## MSE Ridge: 58.75168
```

```
cat('MSE Lasso: ', MSE_lasso, '\n')
```

```
## MSE Lasso: 58.06509
```

```
cat('MSE PCR: ', MSE_pcr, '\n')
```

```
## MSE PCR: 57.61252
```

```
cat('MSE PLS: ', MSE_pls, '\n')
```

```
## MSE PLS: 57.61217
```

We see that PCR and PLS perform the best out of best subset, ridge, lasso, PCR, and PLS, while PLS performs very slightly better than PCR. This is interesting because PCR is just OLS regression in this case since the number of transformed predictors equals the number of original predictors ( $M = p$ ). So we could just use the standard OLS model or use PLS with ten variables if we want a very slight improvement. Because PLS is less interpretable than OLS, we'd probably want to favor using OLS.

(c)

We chose PCR(OLS) as our final model. It includes all of the features in the dataset because that is how PCR works in general, it uses linear combinations of every feature for each transformed predictor. In this case, the PCR model is just the OLS model since  $M = p$ .