

HW6

Jon Griffith and Annabelle Cunningham

2025-04-21

```
library(ISLR2)

## Warning: package 'ISLR2' was built under R version 4.4.3
library(glmnet)

## Warning: package 'glmnet' was built under R version 4.4.3
## Loading required package: Matrix
## Loaded glmnet 4.1-8
library(leaps)

## Warning: package 'leaps' was built under R version 4.4.3
library(pls)

## Warning: package 'pls' was built under R version 4.4.3
##
## Attaching package: 'pls'
## The following object is masked from 'package:stats':
##      loadings
```

8

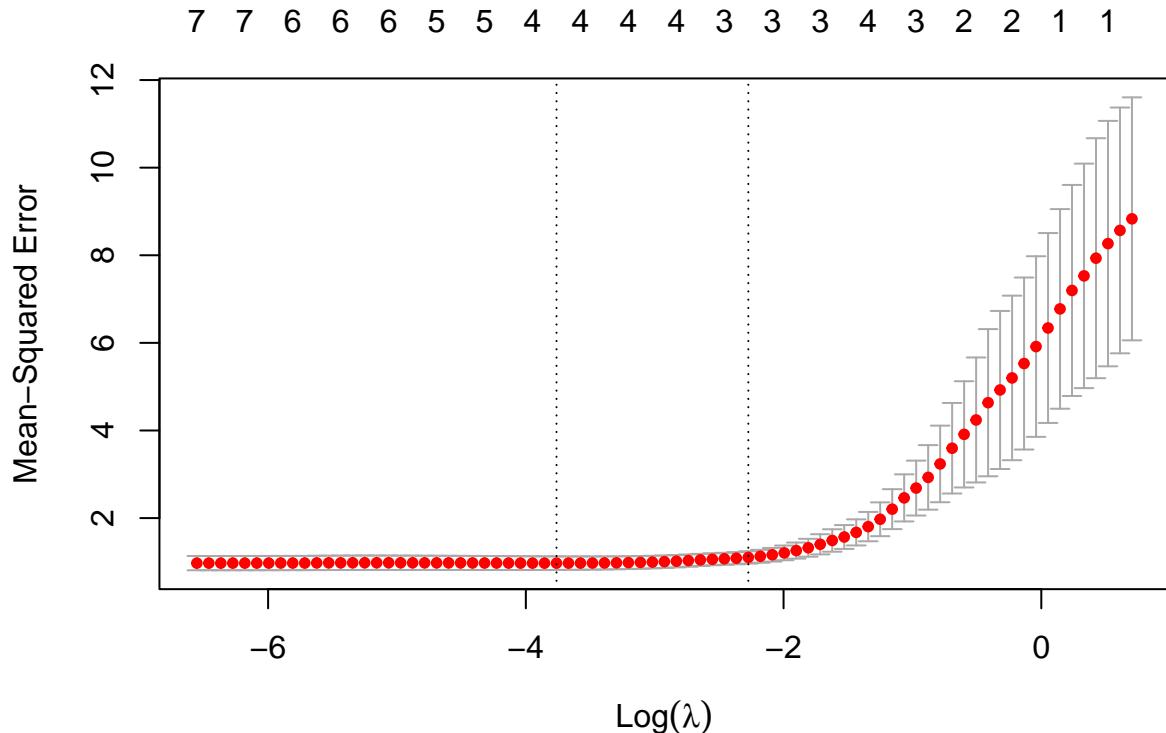
(e)

```
set.seed(1)
# Generate data
x <- rnorm(100)
e <- rnorm(100)
Y <- 4 - 3*x - 2*x^2 + 1.25*x^3 + e

df <- data.frame(x,Y)

set.seed(1)
cv.out <- cv.glmnet(poly(df$x,10,raw=TRUE), df$Y, alpha=1)

plot(cv.out)
```



```
lasso.mod <- glmnet(poly(df$x, 10, raw=TRUE), df$Y, alpha=1, lambda=cv.out$lambda.min)
predict(lasso.mod, type='coefficients', s=cv.out$lambda.min)
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##           s1
## (Intercept) 4.0446480
## 1          -2.3121547
## 2          -2.1331883
## 3           0.5836176
## 4            .
## 5           0.1110500
## 6            .
## 7            .
## 8            .
## 9            .
## 10           .
```

Our resulting coefficient predictions based on the optimal lambda from CV come close to the true values for β_0 , β_1 , and β_2 , while β_3 is about half of its actual value and β_5 is included even though there was no X^5 in the true equation. It still did a pretty good job of selecting close to the actual number of variables in the model and was only off by one.

(f)

```
set.seed(1)
df$Y <- 4 - 1.25*x^7 + e

bestsub.mod <- regsubsets(Y ~ poly(x, 10, raw=TRUE), data=df)

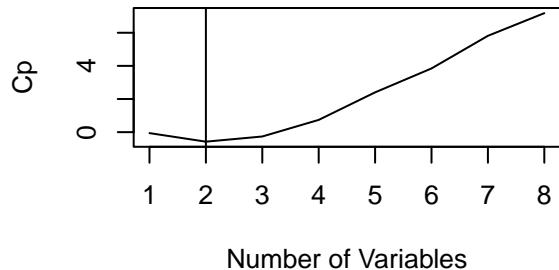
reg.summary <- summary(bestsub.mod)
```

```

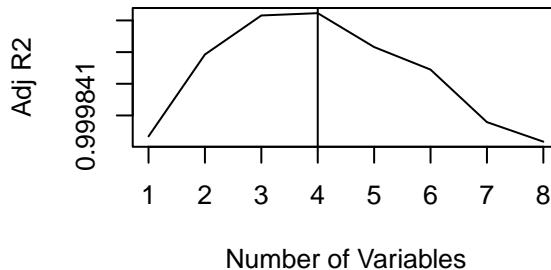
par(mfrow=c(2,2))
# PLOT of cp vs num variables
plot(reg.summary$cp,
  xlab='Number of Variables',
  ylab='Cp',
  type='l',
  main= paste('Optimal:', which.min(reg.summary$cp)))
abline(v=which.min(reg.summary$cp))
# PLOT of adjr2 vs num variables
plot(reg.summary$adjr2,
  xlab='Number of Variables',
  ylab='Adj R2',
  type='l',
  main= paste('Optimal:', which.max(reg.summary$adjr2)))
abline(v=which.max(reg.summary$adjr2))
# PLOT of bic vs num variables
plot(reg.summary$bic,
  xlab='Number of Variables',
  ylab='BIC',
  type='l',
  main= paste('Optimal:', which.min(reg.summary$bic)))
abline(v=which.min(reg.summary$bic))

```

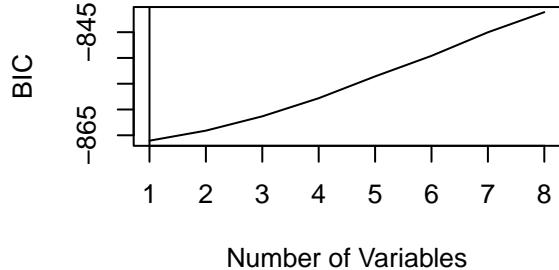
Optimal: 2



Optimal: 4



Optimal: 1



```
coef(bestsub.mod, 1)
```

```
##             (Intercept) poly(x, 10, raw = TRUE)7
##             3.95894          -1.24923
```

```
coef(bestsub.mod, 2)
```

```

##          (Intercept) poly(x, 10, raw = TRUE)2 poly(x, 10, raw = TRUE)7
##          4.0704904           -0.1417084            -1.2484448
coef(bestsub.mod, 4)

##          (Intercept) poly(x, 10, raw = TRUE)1 poly(x, 10, raw = TRUE)2
##          4.0762524           0.2914016           -0.1617671
## poly(x, 10, raw = TRUE)3 poly(x, 10, raw = TRUE)7
##          -0.2526527           -1.2408662

```

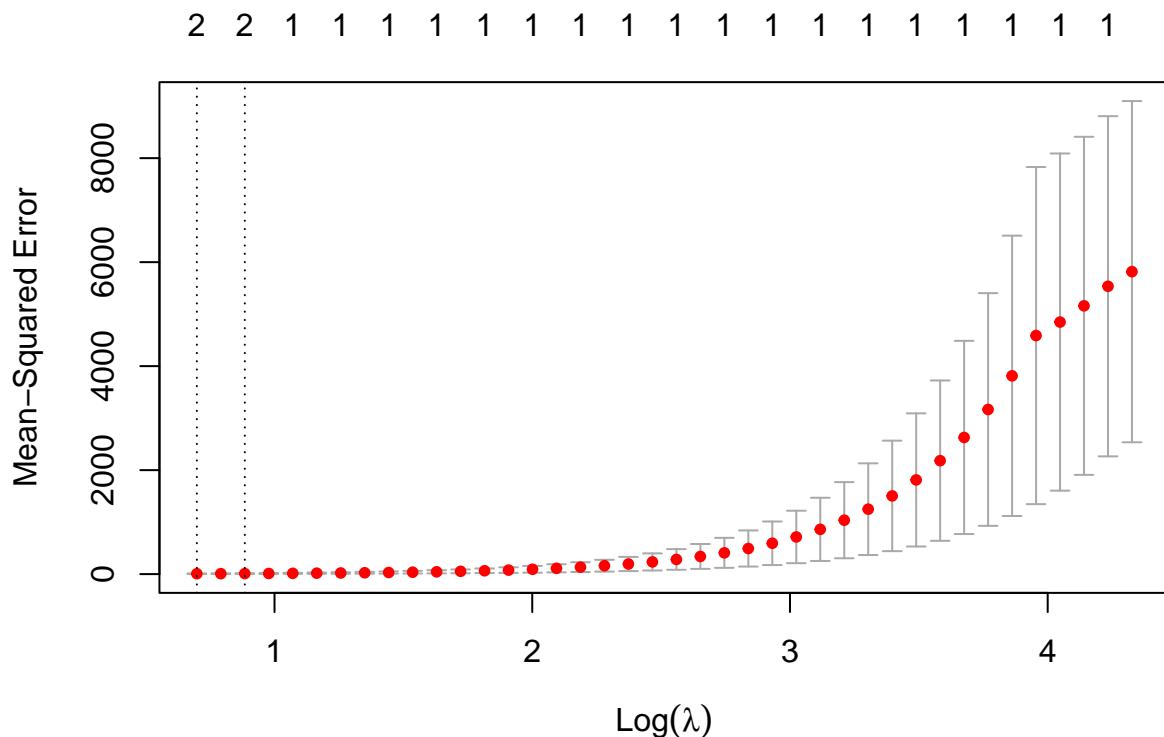
The best models as measured by BIC, CP, and Adj R^2 determined best subsets of (X^7) , (X^2, X^7) , and $(X, X^2, X^3, \text{ and } X^7)$, respectively. Each model included the predictor X^7 that corresponded to the true predictor, and the intercepts and coefficient for X^7 very closely approximated the true values for each. BIC did the best job but CP and Adj R^2 were nearly as good, adding only negligible effects of extra predictors.

```

set.seed(1)
cv.out <- cv.glmnet(poly(df$x, 10, raw=TRUE), df$Y, alpha=1)

plot(cv.out)

```



```

lasso.mod <- glmnet(poly(df$x, 10, raw=TRUE), df$Y, alpha=1, lambda=cv.out$lambda.min)
predict(lasso.mod, type='coefficients', s=cv.out$lambda.min)

```

```

## 11 x 1 sparse Matrix of class "dgCMatrix"
##           s1
## (Intercept) 3.81888390
## 1           .
## 2           .
## 3           .
## 4           .
## 5      -0.01814104
## 6           .
## 7     -1.21252256

```

```
## 8 .
## 9 .
## 10 .
```

The optimal lambda and corresponding model for Lasso, per CV, includes an intercept value of 3.81 which very closely approximates the true intercept of 4, the X^5 predictor with a coefficient of -0.018 , which is not in the true model but is negligible, and the X^7 predictor with a coefficient of -1.212 , which very closely approximates the true coefficient of -1.25 . The Lasso does a good job of replicating the true model and is about on par with the best subsets method.

9

(a)

```
df <- College
head(df)

##                                     Private Apps Accept Enroll Top10perc Top25perc
## Abilene Christian University      Yes 1660    1232    721       23       52
## Adelphi University                Yes 2186    1924    512       16       29
## Adrian College                   Yes 1428    1097    336       22       50
## Agnes Scott College              Yes  417     349    137       60       89
## Alaska Pacific University        Yes  193     146     55       16       44
## Albertson College                Yes  587     479    158       38       62
##                                     F.Undergrad P.Undergrad Outstate Room.Board Books
## Abilene Christian University    2885          537    7440    3300    450
## Adelphi University               2683          1227   12280    6450    750
## Adrian College                  1036           99   11250    3750    400
## Agnes Scott College             510            63   12960    5450    450
## Alaska Pacific University       249            869   7560    4120    800
## Albertson College               678            41  13500    3335    500
##                                     Personal PhD Terminal S.F.Ratio perc.alumni Expend
## Abilene Christian University   2200    70      78    18.1       12    7041
## Adelphi University              1500    29      30    12.2       16   10527
## Adrian College                 1165    53      66    12.9       30    8735
## Agnes Scott College             875     92      97     7.7       37   19016
## Alaska Pacific University      1500    76      72    11.9       2   10922
## Albertson College              675    67      73     9.4       11    9727
##                                     Grad.Rate
## Abilene Christian University    60
## Adelphi University               56
## Adrian College                  54
## Agnes Scott College              59
## Alaska Pacific University       15
## Albertson College                55

set.seed(1)
train <- sample(1:nrow(df), round(0.7*nrow(df)))
```

(b)

```
ols.mod <- lm(Apps ~ ., data=df[train,])
yhat <- predict(ols.mod, newdata = df[-train,])
MSE <- mean((yhat - df$Apps[-train])^2)
```

```

cat('Test MSE for OLS: ', MSE, '\n')

## Test MSE for OLS: 1266407

(c)

set.seed(1)
train_split <- model.matrix(Apps ~ .-1, data=df[train,])
test_split <- model.matrix(Apps ~ .-1, data=df[-train,])
ridge.cv <- cv.glmnet(train_split, df$Apps[train], alpha=0)
yhat <- predict(ridge.cv, newx = test_split, s = ridge.cv$lambda.min)
MSE_ridge <- mean(yhat - df$Apps[-train])^2
cat('Test MSE for Ridge: ', MSE_ridge, '\n')

## Test MSE for Ridge: 1122733
ridge.coef <- predict(ridge.cv, type='coefficients', s=ridge.cv$lambda.min)[1:19,]
ridge.coef

```

	(Intercept)	PrivateNo	PrivateYes	Accept	Enroll
##	-1.835673e+03	2.319012e+02	-2.327310e+02	1.097370e+00	3.870257e-01
##	Top10perc	Top25perc	F.Undergrad	P.Undergrad	Outstate
##	2.656554e+01	2.809031e-02	5.802492e-02	3.525905e-02	-3.238312e-02
##	Room.Board	Books	Personal	PhD	Terminal
##	2.182992e-01	3.007454e-01	-3.084963e-02	-4.137527e+00	-4.295744e+00
##	S.F.Ratio	perc.alumni	Expend	Grad.Rate	
##	1.473171e+01	-4.595790e+00	6.404545e-02	1.003295e+01	

(d)

```

set.seed(1)
lasso.cv <- cv.glmnet(train_split, df$Apps[train], alpha=1)
yhat <- predict(lasso.cv, newx = test_split, s = lasso.cv$lambda.min)
MSE_lasso <- mean(yhat - df$Apps[-train])^2
lasso.coef <- predict(lasso.cv, type='coefficients', s=lasso.cv$lambda.min)[1:19,]

cat('Test MSE for Lasso: ', MSE_lasso, '\n')

```

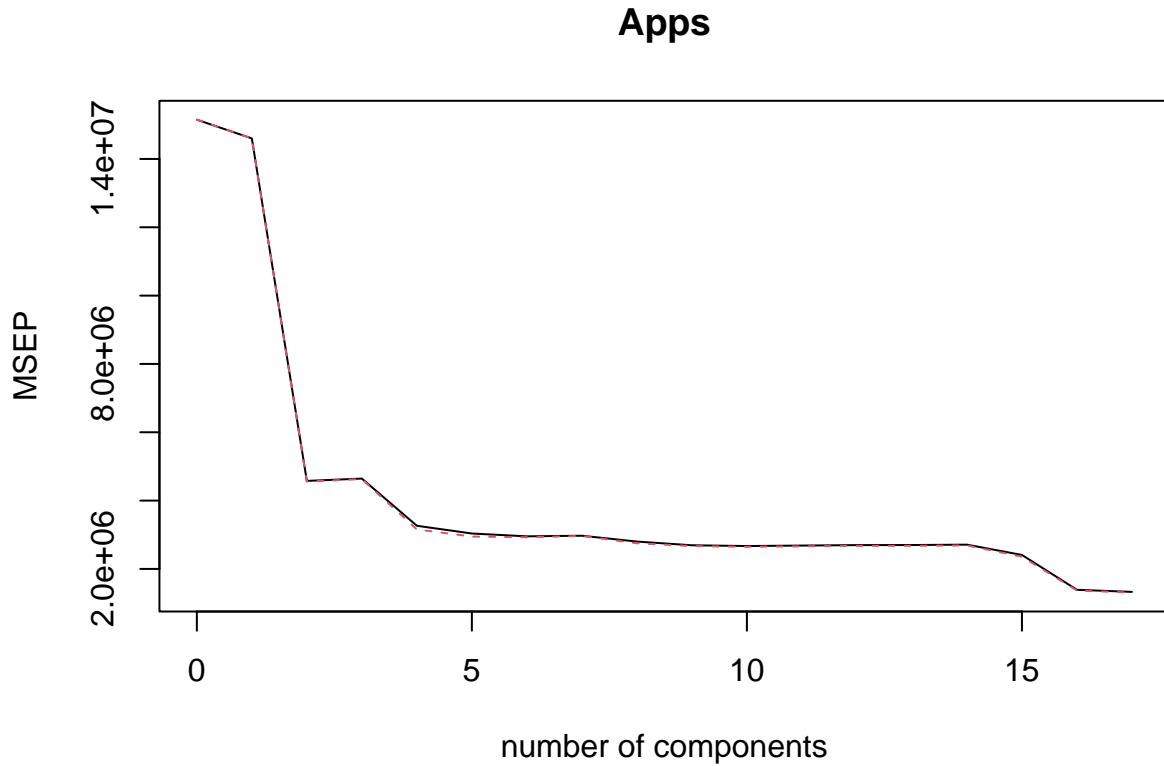
	(Intercept)	PrivateNo	PrivateYes	Accept	Enroll
##	-1.067871e+03	5.064555e+02	-1.271310e-09	1.706280e+00	-9.129561e-01
##	Top10perc	Top25perc	F.Undergrad	P.Undergrad	Outstate
##	5.220892e+01	-1.516972e+01	9.427128e-03	7.078118e-02	-8.550683e-02
##	Room.Board	Books	Personal	PhD	Terminal
##	1.603423e-01	2.622365e-01	0.000000e+00	-9.483740e+00	-8.895288e-02
##	S.F.Ratio	perc.alumni	Expend	Grad.Rate	
##	1.530680e+01	1.523357e+00	5.921301e-02	6.948170e+00	

We have only one predictor zeroed out and several others that are close to zero.

(e)

PCR

```
set.seed(1)
pcr.mod <- pcr(Apps ~ ., data=df[train,], scale=TRUE, validation='CV')
validationplot(pcr.mod, val.type='MSEP')
```



```
which.min(pcr.mod$validation$PRESS)
```

```
## [1] 17
```

We have a minimum at the model with the same number of components as predictors, so we have $M = 17 = p$. This is equivalent to the OLS model, and so we will use this to find our test error equivalent to OLS.

```
M <- 17
yhat <- predict(pcr.mod, df[-train,], ncomp = M)
MSE_pcr <- mean((yhat - df$Apps[-train])^2)

cat('Test MSE for PCR: ', MSE_pcr, '\n')
```

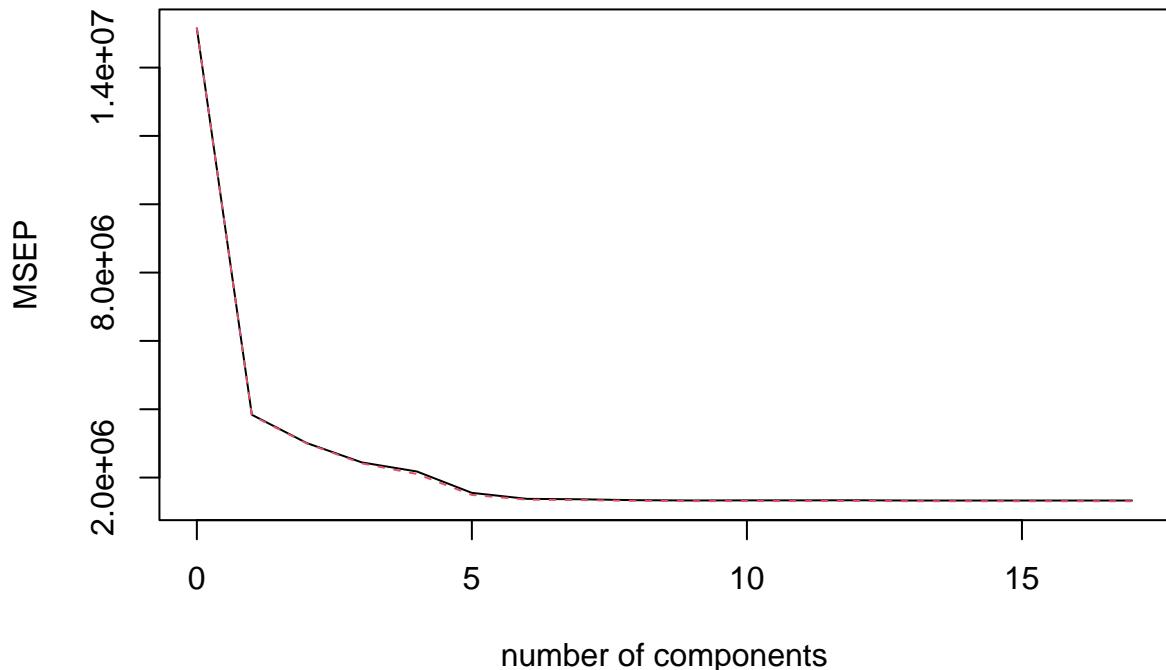
```
## Test MSE for PCR: 1266407
```

(f)

PLS

```
set.seed(1)
pls.mod <- plsr(Apps ~ ., data=df[train,], scale=TRUE, validation='CV')
validationplot(pls.mod, val.type='MSEP')
```

Apps



```
which.min(pls.mod$validation$PRESS)
```

```
## [1] 17
```

We again have a minimum at $M = 17$, which is equivalent to the OLS model, and so we will use this to find our test error equivalent to OLS.

```
M <- 17
yhat <- predict(pls.mod, df[-train,], ncomp = M)
MSE_pls <- mean((yhat - df$Apps[-train])^2)

cat('Test MSE for PLS: ', MSE_pls, '\n')
```

```
## Test MSE for PLS: 1266407
```

(f)

```
cat('MSE OLS: ', MSE, '\n')
```

```
## MSE OLS: 1266407
```

```
cat('MSE Ridge: ', MSE_ridge, '\n')
```

```
## MSE Ridge: 1122733
```

```
cat('MSE Lasso: ', MSE_lasso, '\n')
```

```
## MSE Lasso: 1259544
```

```
cat('MSE PCR: ', MSE_pcr, '\n')
```

```
## MSE PCR: 1266407
```

```
cat('MSE PLS: ', MSE_pls, '\n')
```

```
## MSE PLS: 1266407
```

We achieved our best results from Ridge and Lasso while PCR and PLS just gave the OLS estimate. Out of Ridge and Lasso, Ridge did better.

11

(a)

Create train/test splits

```
df <- Boston

set.seed(1)
train <- sample(1:nrow(df), round(0.7*nrow(df)))

train_split <- model.matrix(crim ~ .-1, data=df[train,])
test_split <- model.matrix(crim ~ .-1, data=df[-train,])
```

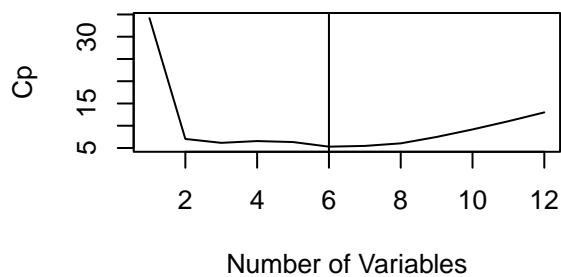
Best Subset Selection

```
bestsub.mod <- regsubsets(crim ~ ., data=df[train,], nvmax=12)

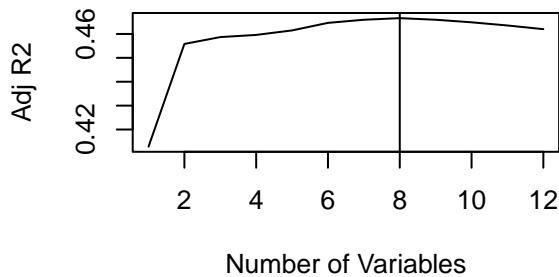
reg.summary <- summary(bestsub.mod)

par(mfrow=c(2,2))
# Plot of cp vs num variables
plot(reg.summary$cp,
  xlab='Number of Variables',
  ylab='Cp',
  type='l',
  main=paste('Optimal:', which.min(reg.summary$cp)))
abline(v=which.min(reg.summary$cp))
# Plot of adjr2 vs num variables
plot(reg.summary$adjr2,
  xlab='Number of Variables',
  ylab='Adj R2',
  type='l',
  main=paste('Optimal:', which.max(reg.summary$adjr2)))
abline(v=which.max(reg.summary$adjr2))
# Plot of bic vs num variables
plot(reg.summary$bic,
  xlab='Number of Variables',
  ylab='BIC',
  type='l',
  main=paste('Optimal:', which.min(reg.summary$bic)))
abline(v=which.min(reg.summary$bic))
```

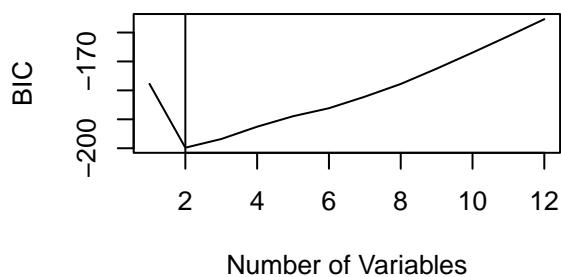
Optimal: 6



Optimal: 8



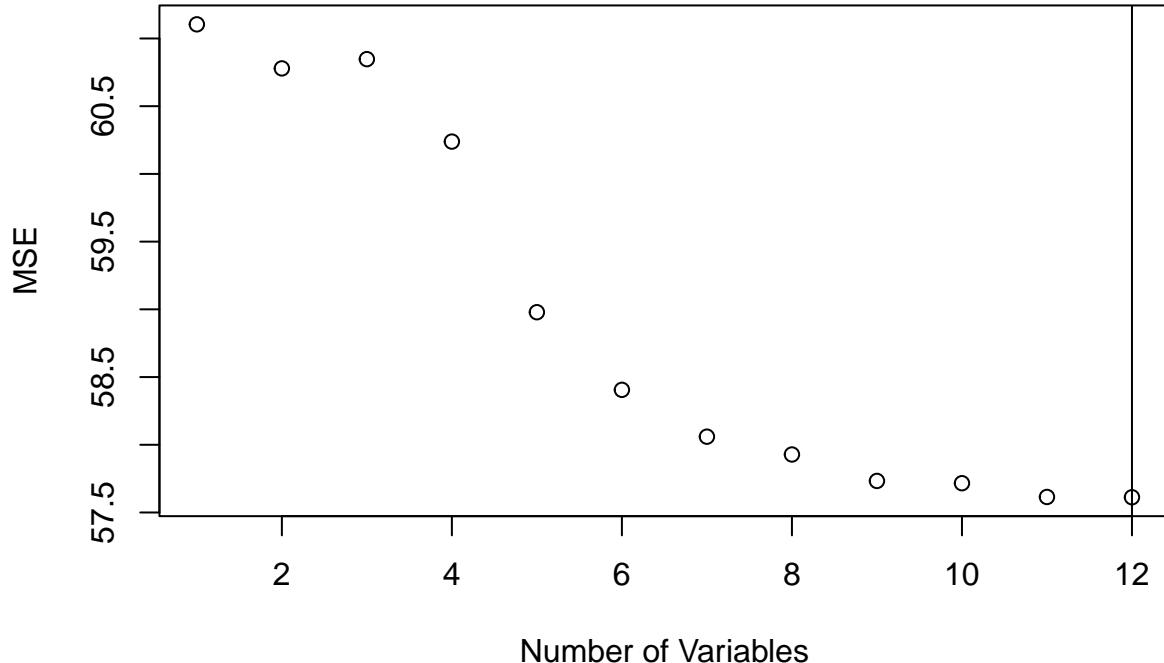
Optimal: 2



```
test.mat <- model.matrix(crim ~ ., data=df[-train,])
val.errors <- rep(NA, 12)
for (i in 1:12){
  coefi <- coef(bestsub.mod, id=i)
  pred <- test.mat[,names(coefi)] %*% coefi
  val.errors[i] <- mean((df$crim[-train] - pred)^2)
}
plot(val.errors,
  main='MSE for each model fit\n Optimal: 7',
  xlab='Number of Variables',
  ylab='MSE')
abline(v = which.min(val.errors))
```

MSE for each model fit

Optimal: 7



```
MSE_bestsub <- val.errors[2]
```

```
MSE_bestsub
```

```
## [1] 60.77902
```

Ridge Regression

```
ridge.cv <- cv.glmnet(train_split, df$crim[train], alpha=0)
yhat <- predict(ridge.cv, newx = test_split, s = ridge.cv$lambda.min)
MSE_ridge <- mean((yhat - df$crim[-train])^2)
MSE_ridge
```

```
## [1] 58.75168
```

Lasso

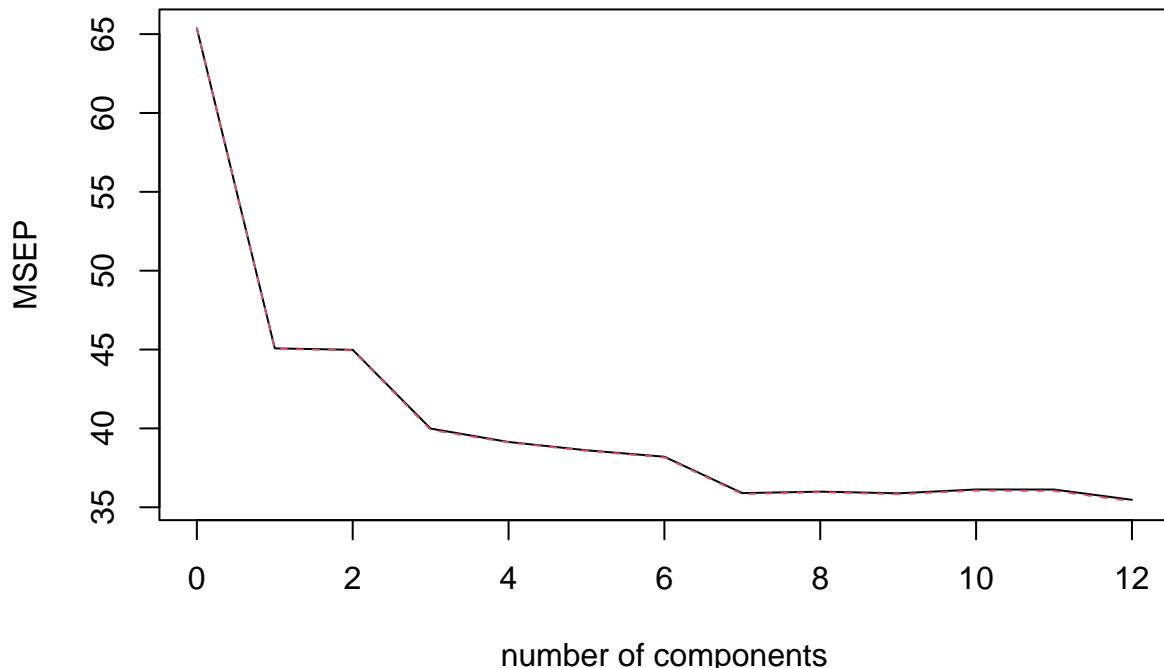
```
lasso.cv <- cv.glmnet(train_split, df$crim[train], alpha=1)
yhat <- predict(lasso.cv, newx = test_split, s = lasso.cv$lambda.min)
MSE_lasso <- mean((yhat - df$crim[-train])^2)
MSE_lasso
```

```
## [1] 58.06509
```

PCR

```
set.seed(1)
pqr.mod <- pqr(crim ~ ., data=df[train,], scale=TRUE, validation='CV')
validationplot(pqr.mod, val.type='MSEP')
```

crim



```
which.min(pcr.mod$validation$PRESS)
```

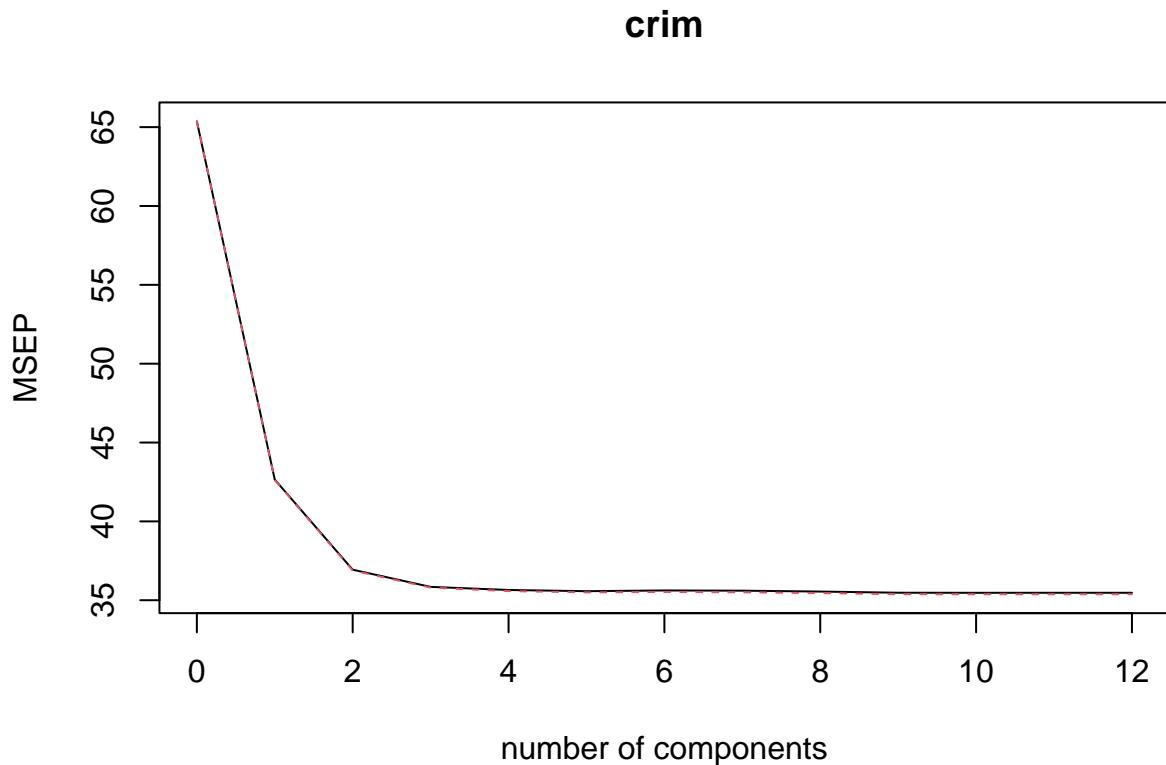
```
## [1] 12  
M <- 12  
yhat <- predict(pcr.mod, df[-train,], ncomp = M)  
MSE_pcr <- mean((yhat - df$crim[-train])^2)
```

```
MSE_pcr
```

```
## [1] 57.61252
```

PLS

```
set.seed(1)  
pls.mod <- plsr(crim ~ ., data=df[train,], scale=TRUE, validation='CV')  
validationplot(pls.mod, val.type='MSEP')
```



```
which.min(pls.mod$validation$PRESS)
```

```
## [1] 10
M <- 10
yhat <- predict(pls.mod, df[-train,], ncomp = M)
MSE_pls <- mean((yhat - df$crim[-train])^2)
```

```
MSE_pls
```

```
## [1] 57.61217
```

(b)

```
cat('MSE Best Subset: ', MSE_bestsub, '\n')
```

```
## MSE Best Subset: 60.77902
```

```
cat('MSE Ridge: ', MSE_ridge, '\n')
```

```
## MSE Ridge: 58.75168
```

```
cat('MSE Lasso: ', MSE_lasso, '\n')
```

```
## MSE Lasso: 58.06509
```

```
cat('MSE PCR: ', MSE_pcr, '\n')
```

```
## MSE PCR: 57.61252
```

```
cat('MSE PLS: ', MSE_pls, '\n')
```

```
## MSE PLS: 57.61217
```

We see that PCR and PLS perform the best out of best subset, ridge, lasso, PCR, and PLS, while PLS performs very slightly better than PCR. This is interesting because PCR is just OLS regression in this case since the number of transformed predictors equals the number of original predictors ($M = p$). So we could just use the standard OLS model or use PLS with ten variables if we want a very slight improvement. Because PLS is less interpretable than OLS, we'd probably want to favor using OLS.

(c)

We chose PCR(OLS) as our final model. It includes all of the features in the dataset because that is how PCR works in general, it uses linear combinations of every feature for each transformed predictor. In this case, the PCR model is just the OLS model since $M = p$.