

# HW7

Jon Griffith and Beatrice Lowe

2025-04-28

```
library(ISLR2)

## Warning: package 'ISLR2' was built under R version 4.4.3
library(boot)
```

## 3

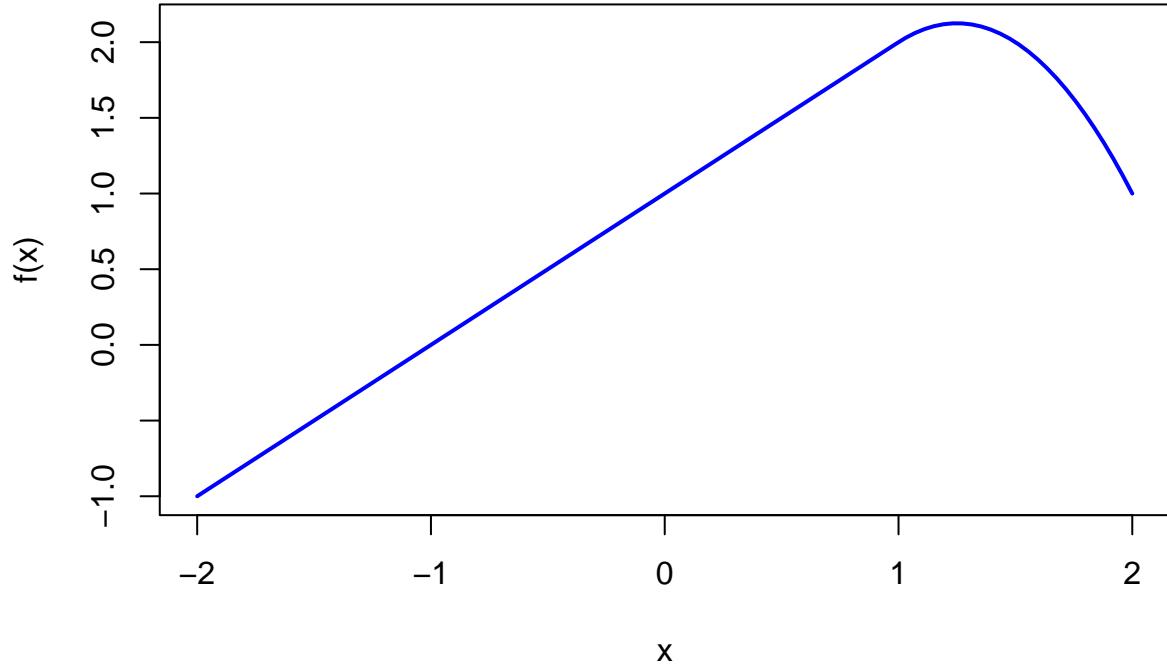
```
x <- seq(-2, 2, length.out=100)
B0 <- 1
B1 <- 1
B2 <- -2

b1 <- x
b2 <- (x-1)^2*I(x >= 1)

f <- function(x) {B0 + B1*b1 + B2*b2}

plot(x, f(x), type='l', lwd=2, col='blue',
      main = expression(f(x) == beta[0] + beta[1]*b[1](x) + beta[2]*b[2](x) + epsilon))
```

$$f(x) = \beta_0 + \beta_1 b_1(x) + \beta_2 b_2(x) + \varepsilon$$



4

```

x <- seq(-2, 6, length.out=500)
B0 <- 1
B1 <- 1
B2 <- 3

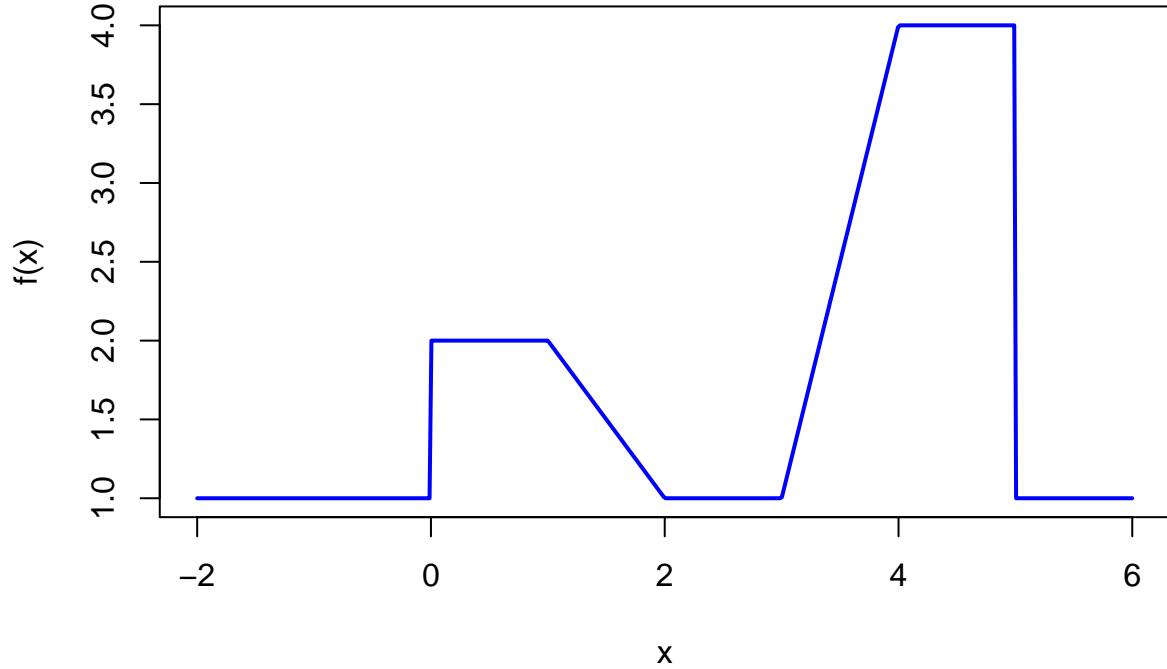
b1 <- (I(0 <= x & x <= 2)-(x-1)*I(1 <= x & x <=2))
b2 <- (x-3)*I(3 <= x & x <= 4) + I(4 < x & x <= 5)

f <- function(x) {B0 + B1*b1 + B2*b2}

plot(x, f(x), type='l', lwd=2, col='blue',
      main = expression(f(x) == beta[0] + beta[1]*b[1](x) + beta[2]*b[2](x) + epsilon))

```

$$f(x) = \beta_0 + \beta_1 b_1(x) + \beta_2 b_2(x) + \varepsilon$$



6

(a)

```

set.seed(1)

cv.error.10 <- rep(NA, 10)
glm.fits <- vector('list', 10)

for (i in 1:10){
  glm.fits[[i]] <- glm(wage ~ poly(age, i), data=Wage)
  cv.error.10[i] <- cv.glm(Wage, glm.fits[[i]], K=10)$delta[1]
}
cat('Min MSE Degree',which.min(cv.error.10), '\n')

## Min MSE Degree 9
do.call(anova, c(glm.fits, test='F'))

## Analysis of Deviance Table
##
## Model 1: wage ~ poly(age, i)
## Model 2: wage ~ poly(age, i)
## Model 3: wage ~ poly(age, i)
## Model 4: wage ~ poly(age, i)
## Model 5: wage ~ poly(age, i)
## Model 6: wage ~ poly(age, i)
## Model 7: wage ~ poly(age, i)
## Model 8: wage ~ poly(age, i)

```

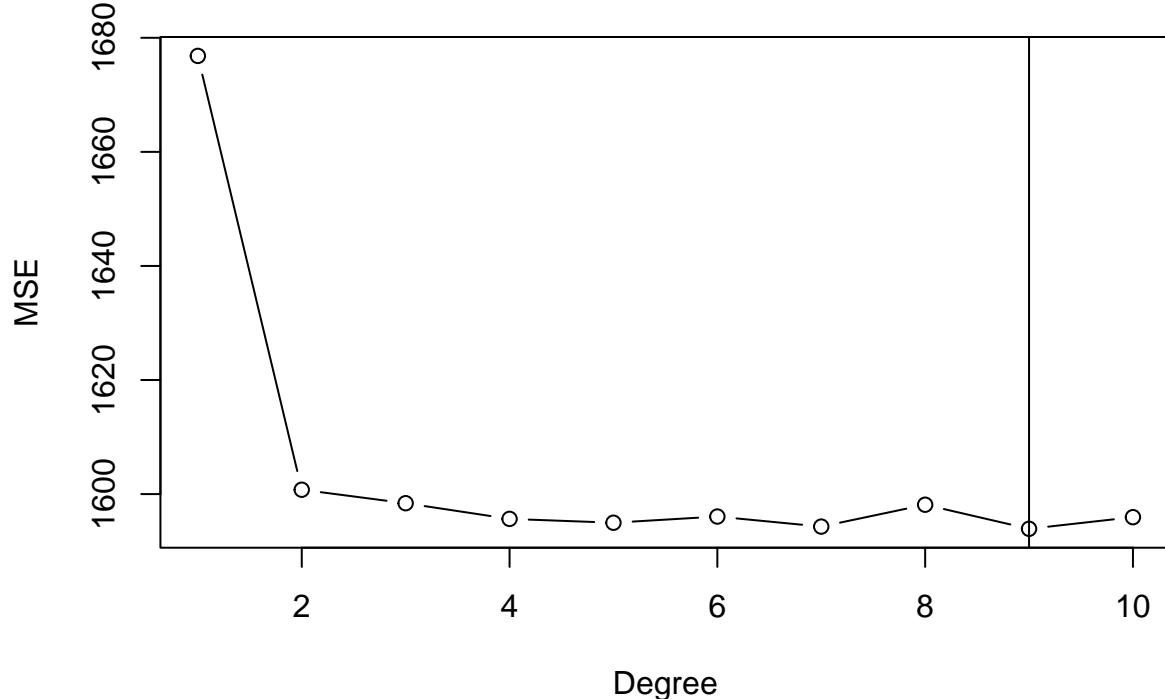
```

## Model 9: wage ~ poly(age, 1)
## Model 10: wage ~ poly(age, 1)
##      Resid. Df Resid. Dev Df Deviance          F    Pr(>F)
## 1      2998   5022216
## 2      2997   4793430  1   228786 143.7638 < 2.2e-16 ***
## 3      2996   4777674  1   15756   9.9005  0.001669 **
## 4      2995   4771604  1   6070   3.8143  0.050909 .
## 5      2994   4770322  1   1283   0.8059  0.369398
## 6      2993   4766389  1   3932   2.4709  0.116074
## 7      2992   4763834  1   2555   1.6057  0.205199
## 8      2991   4763707  1   127    0.0796  0.777865
## 9      2990   4756703  1   7004   4.4014  0.035994 *
## 10     2989   4756701  1       3   0.0017  0.967529
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

plot(1:10, cv.error.10, type='b',
      main = '10-Fold CV for Degree 1-10 Poly',
      xlab = 'Degree',
      ylab = 'MSE')
abline(v = which.min(cv.error.10))

```

### 10-Fold CV for Degree 1–10 Poly



We see from 10-fold CV that the minimum MSE comes from the degree 9 polynomial, but when we plot all corresponding MSE's to each degree, we see that degree 9 only has a marginal advantage over all but degree 1. Using the ANOVA test, we see that all polynomials above degree three, except for nine, are not statistically significant and we deduce from this result that degree three polynomial is probably the optimal degree. We plot this below.

```

agelims <- range(Wage$age)
age.grid <- seq(from=agelims[1], to=agelims[2])
preds <- predict(glm(wage ~ poly(age, 3), data=Wage), newdata=list(age=age.grid), se=TRUE)

```

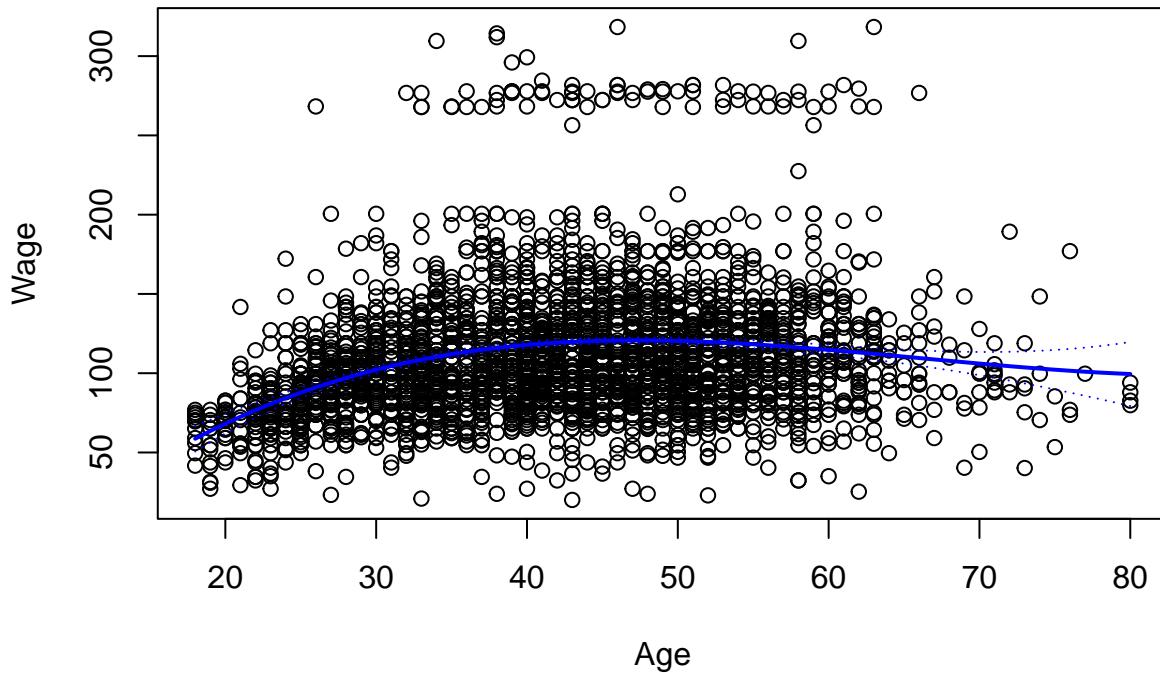
```

se.bands <- cbind(preds$fit + 2*preds$se.fit, preds$fit - 2*preds$se.fit)

plot(Wage$age, Wage$wage,
      main = 'Degree 3 Poly Fit w/ 95% CI',
      xlab = 'Age',
      ylab = 'Wage')
lines(age.grid, preds$fit, lwd=2, col='blue')
matlines(age.grid, se.bands, lwd=1, col='blue', lty=3)

```

**Degree 3 Poly Fit w/ 95% CI**



(b)

```

set.seed(1)

cv.error.10 <- rep(NA, 10)

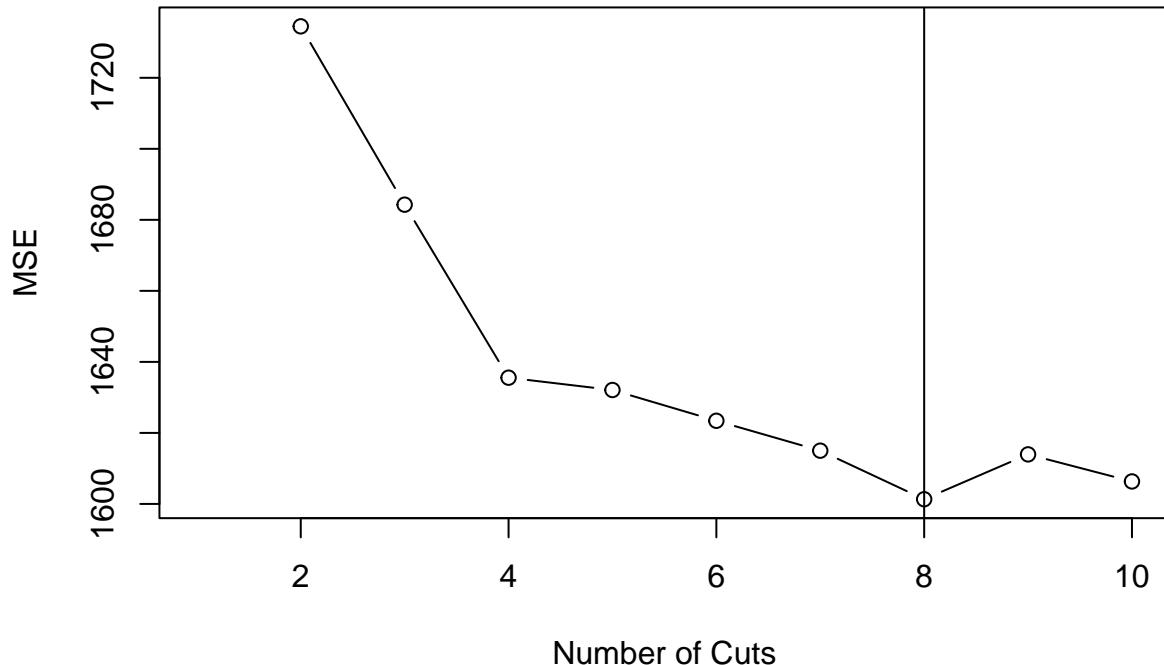
for (i in 2:10){
  Wage$age.cut = cut(Wage$age, i)
  cv.error.10[i] <- cv.glm(wage ~ age.cut, data=Wage, K=10)$delta[1]
}
cat('Min MSE Cut', which.min(cv.error.10), '\n')

## Min MSE Cut 8

plot(1:10, cv.error.10, type='b',
      main = '10-Fold CV for Step Function Cuts 1-10 \n Optimal: 8',
      xlab = 'Number of Cuts',
      ylab = 'MSE')
abline(v = which.min(cv.error.10))

```

## 10–Fold CV for Step Function Cuts 1–10 Optimal: 8

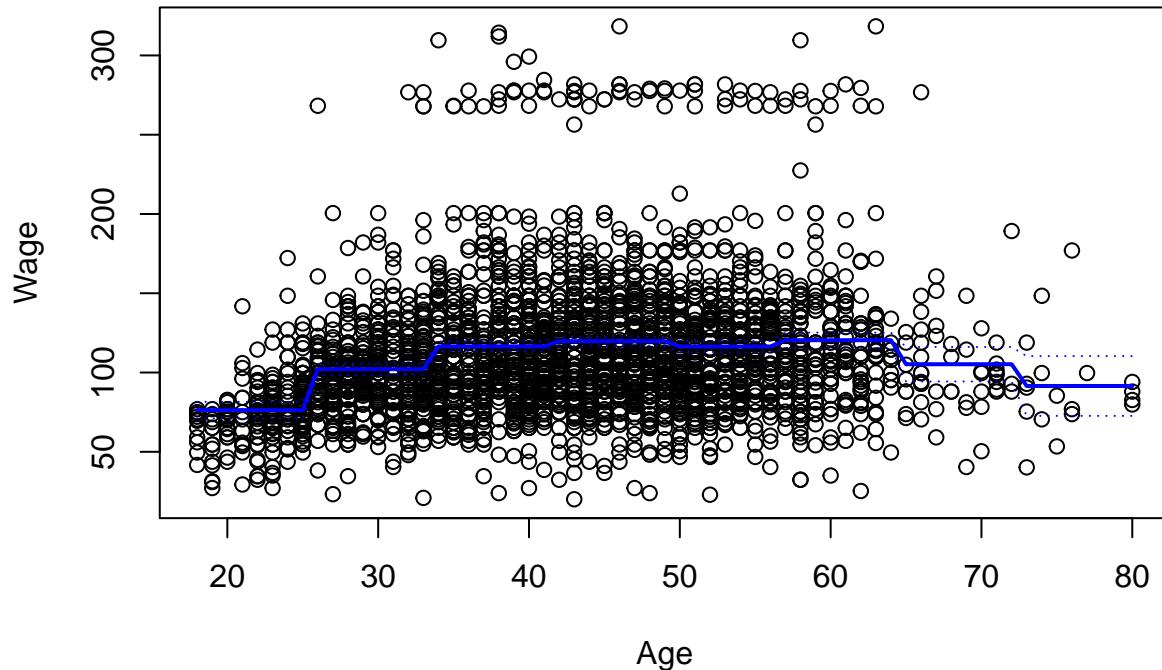


We see from our 10-fold CV results testing cuts 2–10 that the optimal number of cuts based on MSE is eight cuts. We plot this below.

```
agelims <- range(Wage$age)
age.grid <- seq(from=agelims[1], to=agelims[2])
preds <- predict(glm(wage ~ cut(age, 8), data=Wage), newdata=list(age=age.grid), se=TRUE)
se.bands <- cbind(preds$fit + 2*preds$se.fit, preds$fit - 2*preds$se.fit)

plot(Wage$age, Wage$wage,
     main = 'Stepwise Fit w/ 95% CI: 8 Cuts',
     xlab = 'Age',
     ylab = 'Wage')
lines(age.grid, preds$fit, lwd=2, col='blue')
matlines(age.grid, se.bands, lwd=1, col='blue', lty=3)
```

## Stepwise Fit w/ 95% CI: 8 Cuts



**9**

(a)

```
head(Boston)
```

```
##      crim zn indus chas   nox     rm    age     dis rad tax ptratio lstat medv
## 1 0.00632 18 2.31 0 0.538 6.575 65.2 4.0900 1 296 15.3 4.98 24.0
## 2 0.02731 0 7.07 0 0.469 6.421 78.9 4.9671 2 242 17.8 9.14 21.6
## 3 0.02729 0 7.07 0 0.469 7.185 61.1 4.9671 2 242 17.8 4.03 34.7
## 4 0.03237 0 2.18 0 0.458 6.998 45.8 6.0622 3 222 18.7 2.94 33.4
## 5 0.06905 0 2.18 0 0.458 7.147 54.2 6.0622 3 222 18.7 5.33 36.2
## 6 0.02985 0 2.18 0 0.458 6.430 58.7 6.0622 3 222 18.7 5.21 28.7
```

```
poly.fit <- glm(nox ~ poly(dis, 3), data=Boston)
summary(poly.fit)
```

```
##
## Call:
## glm(formula = nox ~ poly(dis, 3), data = Boston)
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 0.554695  0.002759 201.021 < 2e-16 ***
## poly(dis, 3)1 -2.003096  0.062071 -32.271 < 2e-16 ***
## poly(dis, 3)2  0.856330  0.062071 13.796 < 2e-16 ***
## poly(dis, 3)3 -0.318049  0.062071 -5.124 4.27e-07 ***
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```

##  

## (Dispersion parameter for gaussian family taken to be 0.003852802)  

##  

##      Null deviance: 6.7810  on 505  degrees of freedom  

## Residual deviance: 1.9341  on 502  degrees of freedom  

## AIC: -1370.9  

##  

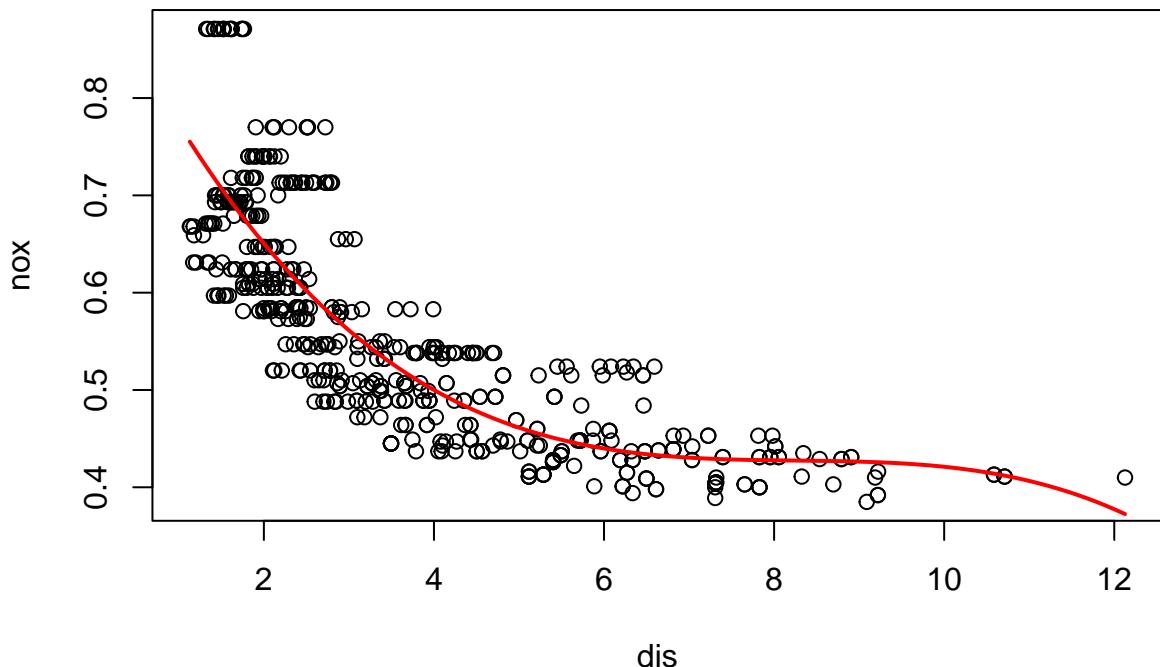
## Number of Fisher Scoring iterations: 2  

dis.range <- range(Boston$dis)
dis.grid <- seq(dis.range[1], dis.range[2], length.out=100)
preds <- predict(poly.fit, newdata=list(dis = dis.grid))

plot(Boston$dis, Boston$nox,
      main='Degree 3 Poly Fit: dis VS nox',
      xlab='dis',
      ylab='nox')
lines(dis.grid, preds, lwd=2, col='red')

```

**Degree 3 Poly Fit: dis VS nox**



(b)

```

dis_range <- range(Boston$dis)
dis.grid <- seq(dis_range[1], dis_range[2], length.out=100)
rss_vec <- rep(NA, 10)
pred_mat <- matrix(NA, nrow=100, ncol=10)

for (i in 1:10){
  fit <- glm(nox ~ poly(dis, i), data=Boston)
  pred_mat[,i] <- predict(fit, newdata=list(dis = dis.grid))
  rss_vec[i] <- sum(fit$residuals^2)
}

```

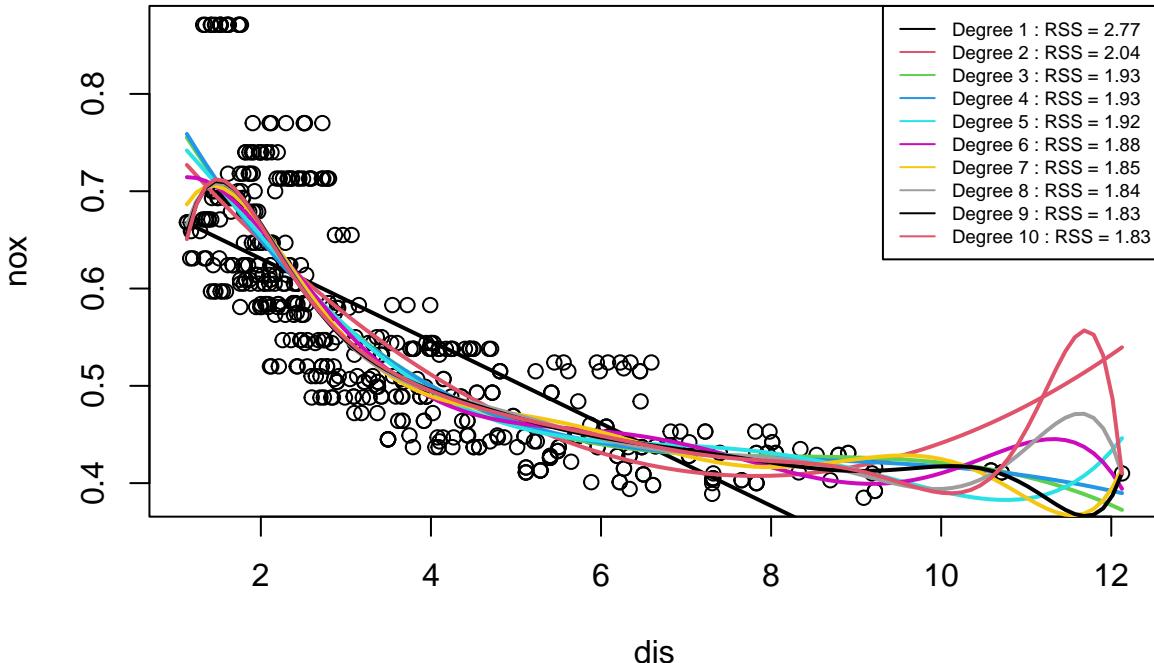
```

}

plot(Boston$dis, Boston$nox,
  main='Degree 1-10 Poly Fits: dis VS nox',
  xlab='dis',
  ylab='nox')
matlines(dis.grid, pred_mat, lty=1, lwd=2, col=c(1:10))
legend('topright', legend=paste('Degree', 1:10, ': RSS =', round(rss_vec[1:10], 2)), col=c(1:10), lty=1, cex=0.8)

```

### Degree 1–10 Poly Fits: dis VS nox



We see in the above plot of 10 different poly fits, from degree one to degree ten, that degree one does poorly but degree two through ten are relatively close. The best degrees are degree nine and degree ten, though it is important to note that tail behavior becomes more extreme at higher order polynomials. This holds true especially for the degree nine and degree ten fit, suggesting overfitting.

(c)

```

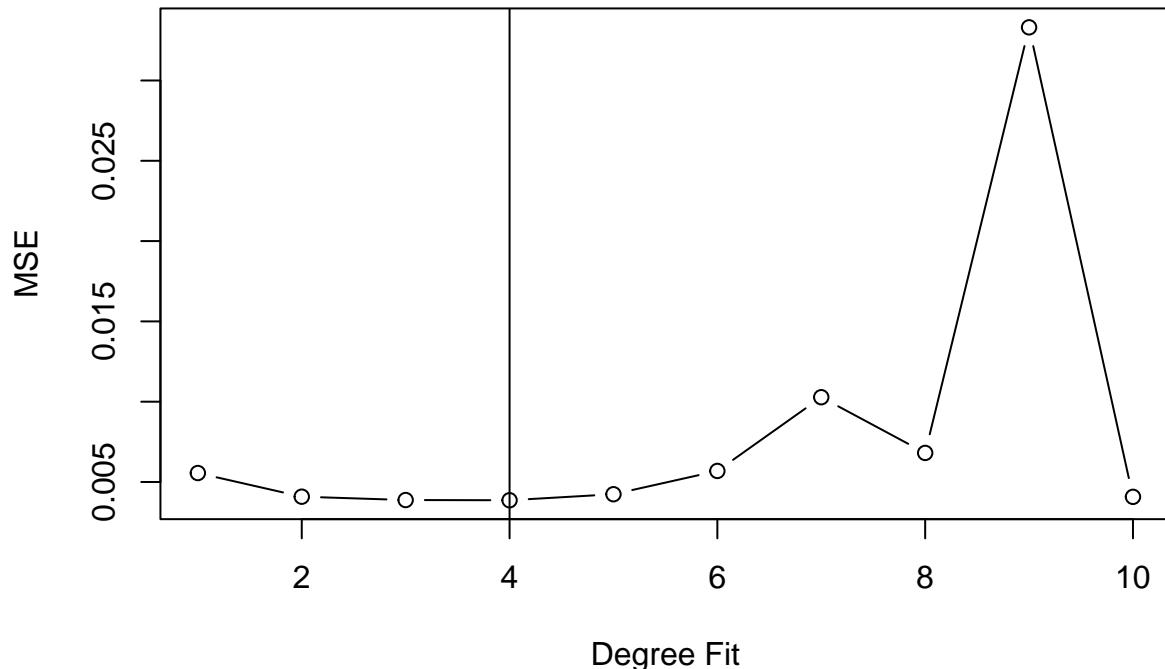
set.seed(1)

cv.error.10 <- rep(NA, 10)
for (i in 1:10){
  fit <- glm(nox ~ poly(dis, i), data=Boston)
  cv.error.10[i] <- cv.glm(Boston, fit, K=10)$delta[1]
}

plot(1:10, cv.error.10, type='b',
  main='Boston: nox VS dis \n MSE for 10-Fold CV on Degree 1-10 Poly fit \n Optimal = 4',
  xlab='Degree Fit',
  ylab='MSE')
abline(v=which.min(cv.error.10))

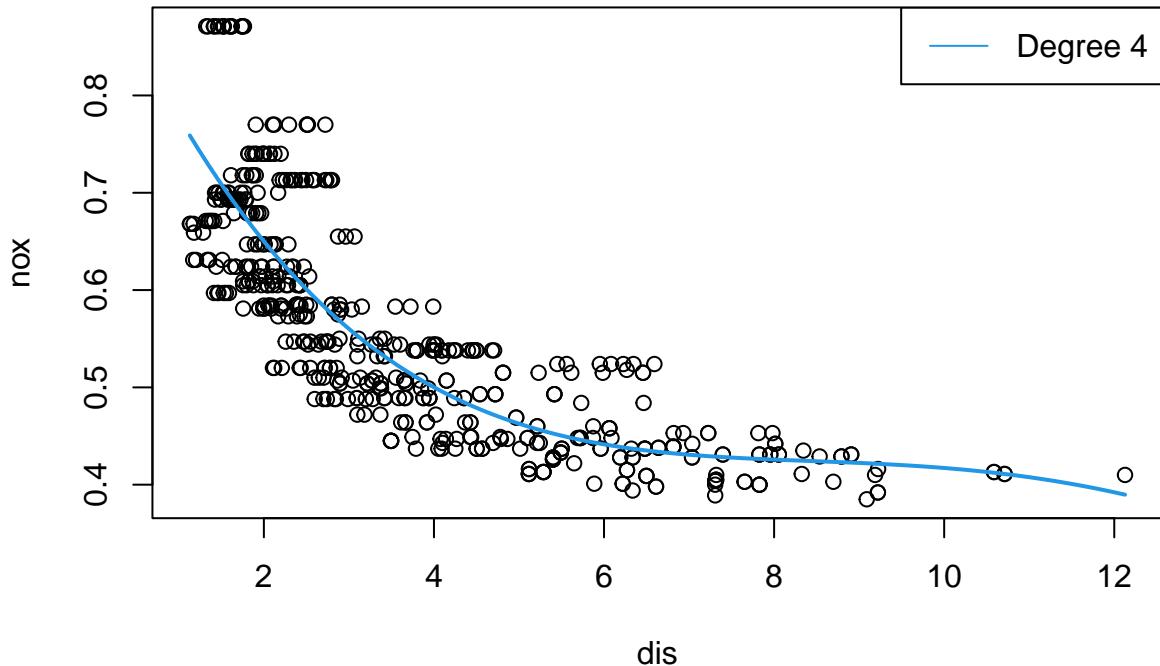
```

**Boston: nox VS dis**  
**MSE for 10–Fold CV on Degree 1–10 Poly fit**  
**Optimal = 4**



```
plot(Boston$dis, Boston$nox,
      main='Degree 4 Poly Fit: dis VS nox',
      xlab='dis',
      ylab='nox')
matlines(dis.grid, pred_mat[,4], lty=1, lwd=2, col=4)
legend('topright', legend=paste('Degree 4'), col=4, lty=1)
```

## Degree 4 Poly Fit: dis VS nox



Using 10-fold CV on degrees one through ten polynomial fits, we obtained four degrees as the optimal fit based on minimizing MSE. It performed relatively well when just fitting the data and recording the RSS, relative to the other fits in the plot in (b). But on that same plot, you can see extreme tail behavior for higher order polynomials which would point to overfitting. Thus, four degrees makes sense since the tail behavior was modest while still giving accurate predictions.